

EEP 523: MOBILE APPLICATIONS FOR SENSING AND CONTROL

SUMMER 2021
Lecture 5

Tamara Bonaci
tbonaci@uw.edu

Agenda

- Canvas - lecture notes
- GitHub repos

HW1
HW2

HW1 - HW6

HW6 becomes optional as well

- HW4, 5 -

7/29/21

Aug 23rd
Aug 22nd - HW5

Review: Sensors in Android

The Android platform supports three broad categories of sensors:

Motion sensors

- Measure: acceleration forces and rotational forces along three axes.
- Types: accelerometers, gravity sensors, gyroscopes, and rotational vector sensors.

Environmental sensors

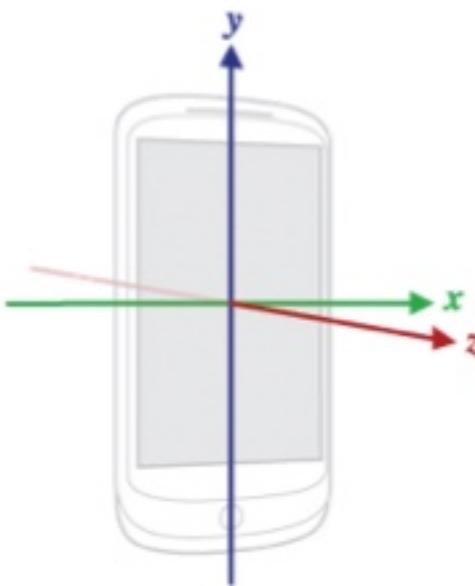
- Measure: ambient air temperature and pressure, illumination, and humidity.
- Types: barometers, photometers, and thermometers.

Position sensors

- Measure: physical position of a device.
- Types: orientation sensors and magnetometers.

Review: Coordinate System

Coordinate system
(relative to a mobile device)
used by the Sensor API.



- The axes are not swapped when the device's screen orientation changes—**the sensor's coordinate system never changes as the device moves**
- your application must not assume that a device's natural (default) orientation is portrait. The natural orientation for many tablet devices is landscape. And **the sensor coordinate system is always based on the natural orientation of a device**.

Review: Accelerometer

- ❑ Acceleration of the device along the 3 sensor axes.
- ❑ The measured acceleration includes both the physical acceleration (change of velocity) and the gravity.
- ❑ The measurement is reported in the x, y and z fields of `sensors_event_t.acceleration`.
- ❑ All values are in SI units (**m/s²**) and measure the acceleration of the device minus the force of gravity along the 3 sensor axes.

Review: Accelerometer in Android

`Sensor.TYPE_ACCELEROMETER`:

<code>SensorEvent.values[0]</code>	Acceleration force along the x axis (including gravity).
<code>SensorEvent.values[1]</code>	Acceleration force along the y axis (including gravity).
<code>SensorEvent.values[2]</code>	Acceleration force along the z axis (including gravity).

The force of gravity is always influencing the measured acceleration

- Device sitting on a table (and not accelerating), the accelerometer reads a magnitude of $g = 9.81 \text{ m/s}^2$
- Device in free fall and therefore rapidly accelerating toward the ground at 9.81 m/s^2 , its accelerometer reads a magnitude of $g = 0 \text{ m/s}^2$

Review: Accelerometer in Android

To measure the real acceleration of the device, the contribution of the force of gravity must be removed from the accelerometer data.

- Use a low-pass filter to isolate the force of gravity
- Apply a high-pass filter to remove the force of gravity

```
override fun onSensorChanged(event: SensorEvent) {  
    // alpha is calculated as t / (t + dT)  
    // with t, the low-pass filter's time-constant  
    // and dT, the event delivery rate  
  
    val alpha = 0.8f  
  
    gravity[0] = alpha * gravity[0] + (1 - alpha) * event.values[0]  
    gravity[1] = alpha * gravity[1] + (1 - alpha) * event.values[1]  
    gravity[2] = alpha * gravity[2] + (1 - alpha) * event.values[2]  
  
    linear_acceleration[0] = event.values[0] - gravity[0]  
    linear_acceleration[1] = event.values[1] - gravity[1]  
    linear_acceleration[2] = event.values[2] - gravity[2]  
}
```

t: rough representation of the latency
that the filter adds to the sensor events,

dt: the sensor's event delivery rate

Review: Accelerometer in Android

- Not necessary to get accelerometer events at a very high rate
- By using a slower rate (SENSOR_DELAY_UI), we get an automatic low-pass filter, which "extracts" the gravity component of the acceleration.
- As an added benefit, we use less power and CPU resources.

```
mSensorManager.registerListener(this, mSensor, SensorManager.SENSOR_DELAY_GAME)
```

Gyroscope

EE P 523, Lecture 5

7/29/21

EEP 523, Summer 2021 - Lecture 5

-DBS - Jeff's work
-Haptic password

Gyroscope

Measure rotational (angular) velocity in 1, 2, or 3 directions

- 3-axis gyroscopes often implemented with a 3-axis accelerometer to provide a 6 degree-of-freedom motion tracking system
- Basic types of gyroscope:
 - Rotary (classical) gyroscopes
 - Vibrating Structure Gyroscope
 - Optical Gyroscopes

Gyroscope

Applications

- The gyroscope helps the accelerometer out with understanding which way your phone is orientated—it adds another level of precision, so those 360-degree photo spheres really look as impressive as possible.
- Gyroscopes aren't exclusive to phones.
They're used in altimeters inside aircraft to determine altitude and position, for example, and to keep cameras steady on the move.

Gyroscope

Gyroscope vs Accelerometer

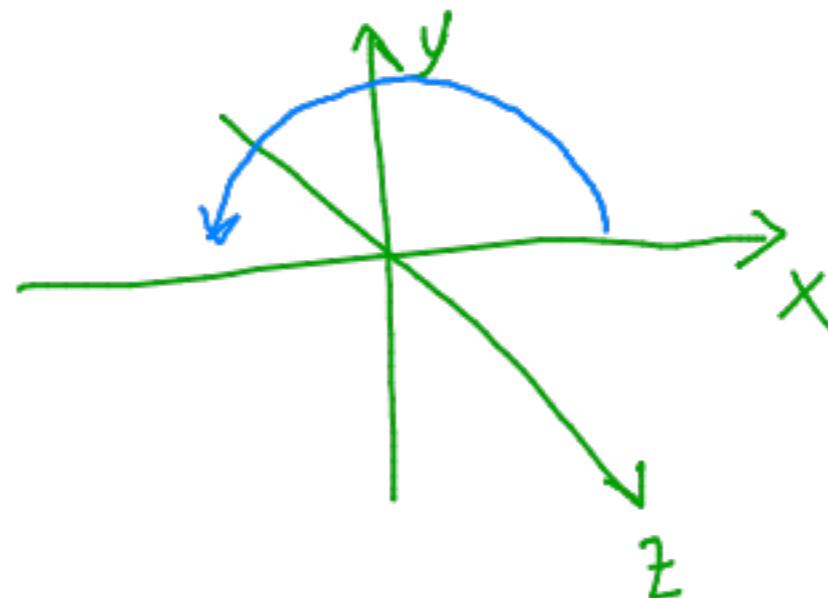
- So, putting in simplest of the words – **gyroscope** provides velocity, and **accelerometer** provides speed.
- Gyroscope is capable of providing precision motion inside the app functionality.
- User can execute majority of the tasks with the motion of the device only.

Gyroscope in Android

```
val sensorManager = getSystemService(Context.SENSOR_SERVICE) as SensorManager  
val sensor: Sensor? = sensorManager.getDefaultSensor(Sensor.TYPE_GYROSCOPE)
```

- All values are in radians/second and measure the rate of rotation around the device's local X, Y and Z axis.
- Rotation is positive in the counter-clockwise direction. That is, an observer looking from some positive location on the x, y or z axis at a device positioned on the origin would report positive rotation if the device appeared to be rotating counter-clockwise.

values[0]: Angular speed around the x-axis
values[1]: Angular speed around the y-axis
values[2]: Angular speed around the z-axis



Gyroscope in Android

```
Override fun onSensorChanged(event: SensorEvent) {
    // This time step's delta rotation to be multiplied by the current rotation
    // after computing it from the gyro sample data.
    if (timestamp != 0) {
        val dT = (event.timestamp - timestamp) * NS2S
        // Axis of the rotation sample, not normalized yet.
        var axisX = event.values[0]
        var axisY = event.values[1]
        var axisZ = event.values[2]
        // Calculate the angular speed of the sample
        val omegaMagnitude = sqrt((axisX * axisX + axisY * axisY + axisZ * axisZ).toDouble())
        // Normalize the rotation vector if it's big enough to get the axis
        if (omegaMagnitude > EPSILON) {
            axisX /= omegaMagnitude.toFloat()
            axisY /= omegaMagnitude.toFloat()
            axisZ /= omegaMagnitude.toFloat()
        }
        // Integrate around this axis with the angular speed by the time step
        // in order to get a delta rotation from this sample over the time step
        // We will convert this axis-angle representation of the delta rotation
        // into a quaternion before turning it into the rotation matrix.
        val thetaOverTwo = omegaMagnitude * dT / 2.0f
        val sinThetaOverTwo = sin(thetaOverTwo)
        val cosThetaOverTwo = cos(thetaOverTwo)
        deltaRotationVector[0] = (sinThetaOverTwo * axisX).toFloat()
        deltaRotationVector[1] = (sinThetaOverTwo * axisY).toFloat()
        deltaRotationVector[2] = (sinThetaOverTwo * axisZ).toFloat()
        deltaRotationVector[3] = cosThetaOverTwo.toFloat()
    }
    timestamp = event.timestamp.toInt()
    val deltaRotationMatrix = FloatArray(9)
    SensorManager.getRotationMatrixFromVector(deltaRotationMatrix, deltaRotationVector)
    // User code should concatenate the delta rotation we computed with the current
    // rotation in order to get the updated rotation.
    // rotationCurrent = rotationCurrent * deltaRotationMatrix;
}
```

Typically, the output of the gyroscope is integrated over time to calculate a rotation describing the change of angles over the time step

Additional Base sensors

- **Compass**
- **Proximity sensor:**
When you lift phone to your ear, the proximity sensor can turn off the display to save power and prevent accidental dialing.
- **Ambient light sensor:**
Automatically brightens the display when you're in sunlight or a bright room and dims it in darker places.

Composite Sensors

EE P 523, Lecture 5

Composite Sensors

Generates data by processing and/or fusing data from one or several physical sensors -> Any sensor that is not a base sensor is called a composite sensor

Examples of composite sensors include:

- **Step detector** and **Significant motion**
usually based on an accelerometer, but could be based on other sensors as well, if the power consumption and accuracy was acceptable.
- **Game rotation vector**, based on an accelerometer and a gyroscope.
- **Uncalibrated gyroscope**, similar to the gyroscope base sensor, but with the bias calibration being reported separately instead of being corrected in the measurement.

Composite Sensors

Sensor type	Category	Underlying physical sensors	Reporting mode
Game rotation vector	Attitude	Accelerometer, Gyroscope MUST NOT USE Magnetometer	Continuous
Geomagnetic rotation vector  %	Attitude	Accelerometer, Magnetometer, MUST NOT USE Gyroscope	Continuous
Glance gesture  %	Interaction	Undefined	One-shot
Gravity	Attitude	Accelerometer, Gyroscope	Continuous
Gyroscope uncalibrated	Uncalibrated	Gyroscope	Continuous
Linear acceleration	Activity	Accelerometer, Gyroscope (if present) or Magnetometer (if gyro not present)	Continuous
Magnetic field uncalibrated	Uncalibrated	Magnetometer	Continuous
Orientation (deprecated)	Attitude	Accelerometer, Magnetometer PREFERRED Gyroscope	Continuous
Pick up gesture  %	Interaction	Undefined	One-shot
Rotation vector	Attitude	Accelerometer, Magnetometer, AND (when present) Gyroscope	Continuous
Significant motion  %	Activity	Accelerometer (or another as long as very low power)	One-shot
Step counter  %	Activity	Accelerometer	On-change
Step detector  %	Activity	Accelerometer	Special
Tilt detector  %	Activity	Accelerometer	Special
Wake up gesture  %	Interaction	Undefined	One-shot

Reading Data from Sensors

1. Class must implement interface *SensorEventListener*

```
class MainActivity : AppCompatActivity(), SensorEventListener {
```

1)

2. Override two functions

```
override fun onAccuracyChanged(sensor: Sensor, accuracy: Int) {}
```



```
override fun onSensorChanged(event: SensorEvent) {  
    if (event.sensor.type != Sensor.TYPE_ACCELEROMETER)  
        return  
    ...  
}
```

Reading Data from Sensors

3. Create a **SensorManager** and **Sensor** variable

```
private lateinit var msensorManager: SensorManager
```

```
private val mAccelerometer: Sensor
```

4. Initiate an instance of the SensorManager and Sensor

```
sensorManager = getSystemService(Context.SENSOR_SERVICE) as SensorManager
```

```
mAccelerometer = mSensorManager!!.getDefaultSensor(Sensor.TYPE_ACCELEROMETER)
```

5. Register a listener for the particular sensor

```
msensorManager!!.registerListener(this, mAccelerometer, SensorManager.SENSOR_DELAY_GAME)
```

Sensor Best Practices

Be sure to unregister a sensor's listener when you are done using the sensor or when the sensor activity pauses.

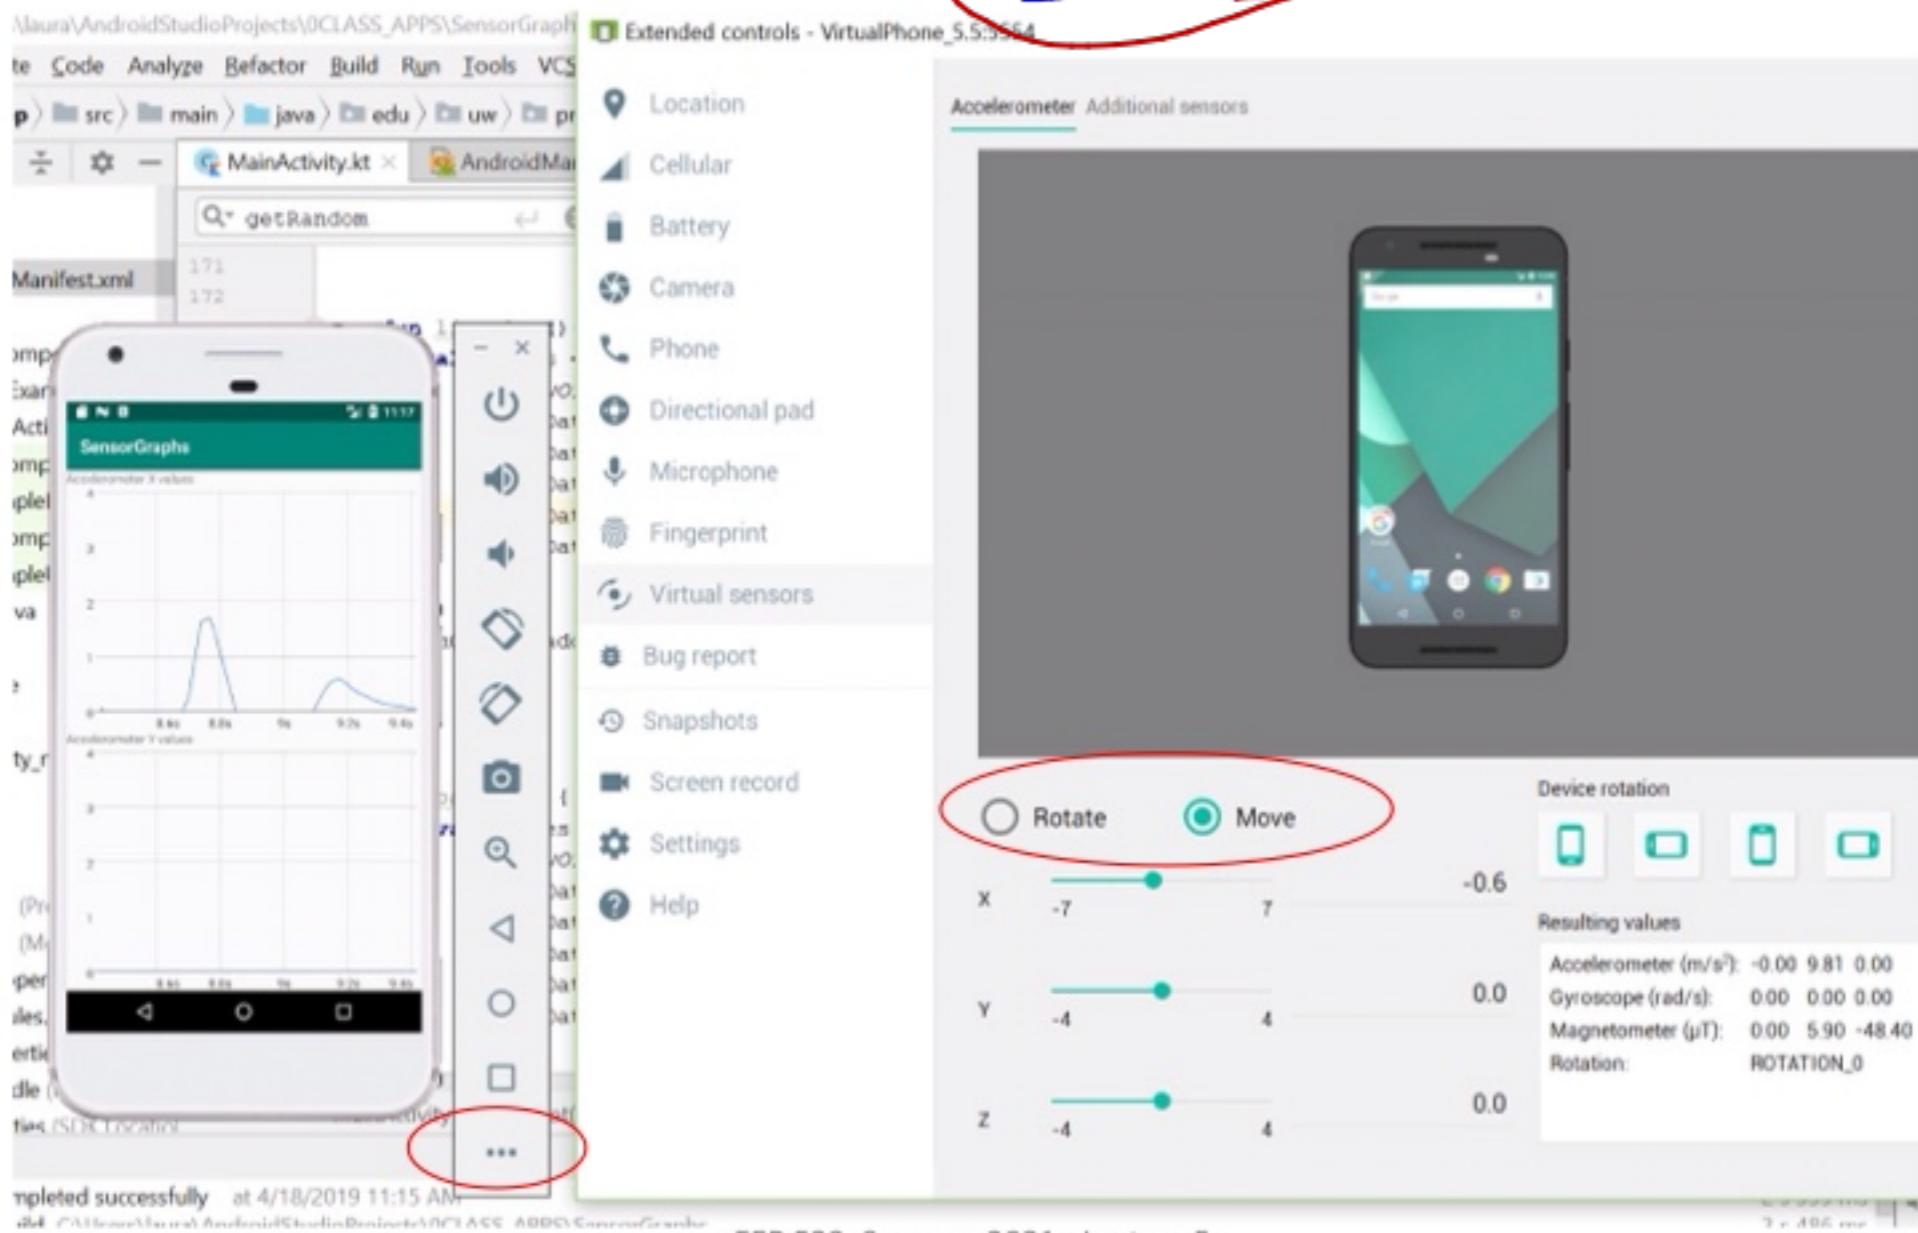
If a listener is registered and its activity is paused, the sensor will continue to acquire data and use the battery resources

```
override fun onPause() {  
    super.onDestroy()  
    mSensorManager.unregisterListener(this)  
}
```

Sensor Best Practices

- Only gather sensor data in the foreground
- On devices running Android 9 (API level 28) or higher, apps running in the background have the following restrictions:
 - Sensors that use the continuous reporting mode ([link](#)), such as accelerometers and gyroscopes, don't receive events.
 - Sensors that use the on-change or one-shot reporting modes don't receive events.
- Given these restrictions, it's best to detect sensor events either when your app is in the foreground or as part of a foreground service.

Sensors in Emulator



7/29/21

EEP 523, Summer 2021 - Lecture 5

Android Supported Sensors

Sensor	Type	Description	Common Uses
<u>TYPE_ACCELEROMETER</u>	Hardware	Measures the acceleration force in m/s^2 that is applied to a device on all three physical axes (x, y, and z), including the force of gravity.	Motion detection (shake, tilt, etc.).
<u>TYPE_AMBIENT_TEMPERATURE</u>	Hardware	Measures the ambient room temperature in degrees Celsius ($^{\circ}\text{C}$). See note below.	Monitoring air temperatures.
<u>TYPE_GRAVITY</u>	Software or Hardware	Measures the force of gravity in m/s^2 that is applied to a device on all three physical axes (x, y, z).	Motion detection (shake, tilt, etc.).
<u>TYPE_GYROSCOPE</u>	Hardware	Measures a device's rate of rotation in rad/s around each of the three physical axes (x, y, and z).	Rotation detection (spin, turn, etc.).
<u>TYPE_LIGHT</u>	Hardware	Measures the ambient light level (illumination) in lx.	Controlling screen brightness.
<u>TYPE_LINEAR_ACCELERATION</u>	Software or Hardware	Measures the acceleration force in m/s^2 that is applied to a device on all three physical axes (x, y, and z), excluding the force of gravity.	Monitoring acceleration along a single axis.

Android Supported Sensors

<u>TYPE_MAGNETIC_FIELD</u>	Hardware	Measures the ambient geomagnetic field for all three physical axes (x, y, z) in μT .	Creating a compass.
<u>TYPE_ORIENTATION</u>	Software	Measures degrees of rotation that a device makes around all three physical axes (x, y, z). As of API level 3 you can obtain the inclination matrix and rotation matrix for a device by using the gravity sensor and the geomagnetic field sensor in conjunction with the getRotationMatrix() method.	Determining device position.
<u>TYPE_PRESSURE</u>	Hardware	Measures the ambient air pressure in hPa or mbar.	Monitoring air pressure changes.
<u>TYPE_PROXIMITY</u>	Hardware	Measures the proximity of an object in cm relative to the view screen of a device. This sensor is typically used to determine whether a handset is being held up to a person's ear.	Phone position during a call.

Android Supported Sensors

<u>TYPE_RELATIVE_HUMIDITY</u>	Hardware	Measures the relative ambient humidity in percent (%).	Monitoring dewpoint, absolute, and relative humidity.
<u>TYPE_ROTATION_VECTOR</u>	Software or Hardware	Measures the orientation of a device by providing the three elements of the device's rotation vector.	Motion detection and rotation detection.
<u>TYPE_TEMPERATURE</u>	Hardware	Measures the temperature of the device in degrees Celsius (°C). This sensor implementation varies across devices and this sensor was replaced with the TYPE_AMBIENT_TEMPERATURE sensor in API Level 14	Monitoring temperatures.

Libraries in Android

EE P 523, Lecture 5

Libraries

- Many Android developers have produced useful **libraries**.
- There is a Maven repository to store various libraries.
- This makes it easy to add them to your Android Studio projects.
- Most libraries use permissive licenses so that you can use them for free and can include them in the code of commercial apps/products.
- (Some libraries must be downloaded as .JARs and added manually to your project)

Adding a Library to Your Project

- Edit the **build.gradle** file for your 'app' module and add lines to the following section at the bottom.
- You can usually find out what file name to write below by going to various libraries' home pages / GitHub pages.



The screenshot shows the Android Studio interface with the project 'BarGraph' open. The left sidebar displays the project structure under 'Resource Manager'. The 'Gradle Scripts' section is expanded, showing the 'build.gradle' file for the 'app' module selected. The code editor on the right contains the Gradle build script. A blue oval highlights the 'dependencies' block, and another blue oval highlights the specific line 'implementation 'com.jjoe64:graphview:4.2.2''.

```
dependencies {
    implementation fileTree(dir: 'libs', include: ['*.jar'])
    implementation "org.jetbrains.kotlin:kotlin-stdlib-jdk7:$kotlin_version"
    implementation 'com.android.support:appcompat-v7:28.0.0'
    implementation 'com.android.support.constraint:constraint-layout:1.1.3'
    implementation 'com.android.support:support-v4:28.0.0'
    testImplementation 'junit:junit:4.12'
    androidTestImplementation 'com.android.support.test:runner:1.0.2'
    androidTestImplementation 'com.android.support.test.espresso:espresso-core:3.0.2'
    implementation 'com.jjoe64:graphview:4.2.2'
}
```

Picasso Library

- **Picasso** is a powerful library for manipulating images.
 - written by Square, inc.
 - <http://square.github.io/picasso/>
- To add Picasso to your project:

```
MAVEN
[<dependency>
 <groupId>com.squareup.picasso</groupId>
 <artifactId>picasso</artifactId>
 <version>(insert latest version)</version>
</dependency>

GRADLE
implementation 'com.squareup.picasso:picasso:(insert latest version)'
```



Picasso Library

In your app's Kotlin code, write:

```
Picasso.with(this).load("url").into(ImageView)
```

Example

```
ImageView img = (ImageView) findViewById(R.id.photo);
Picasso.with(this)
    .load("http://i.imgur.com/DvpvklR.png")
    .into(img);
```

Picasso Library

Method	Description
<code>centerCrop()</code>	center and crop image inside view
<code>centerInside()</code>	resize image proportionally inside view
<code>error(<i>id</i>)</code>	show given drawable as error
<code>fetch()</code>	download image in the background
<code>fit()</code>	resize image to fit view bounds
<code>get()</code>	return image as a Bitmap
<code>into(<i>view</i>)</code>	puts image into given view
<code>placeholder(<i>id</i>)</code>	show given drawable while loading
<code>resize(<i>width, height</i>)</code>	change image size in pixels
<code>rotate(<i>degrees</i>)</code>	rotate clockwise
<code>tag("tag")</code>	attaches a "tag" to a loading image (useful for bulk operations shown later)
<code>transform(<i>trans</i>)</code>	apply complex transformations

Picasso Library

Method	Description
<code>cancelRequest(view)</code>	abort any image loading in that view
<code>cancelTag("tag")</code>	cancel all images with given tag
<code>invalidate("url")</code>	flush out cache of given image,
<code>invalidate(File)</code>	so it will be re-downloaded the next time
<code>load("url")</code>	load an image from various sources
<code>load(id)</code>	
<code>load(File)</code>	
<code>pauseTag("tag")</code>	pause all image loads for given tag
<code>resumeTag("tag")</code>	unpause all image loads for given tag
<code>shutdown()</code>	stop entire Picasso system
<code>with(context)</code>	use given activity/fragment as context

Libraries of Interest

- **Android-Bootstrap**

Customizable widgets not normally available in Android
<https://github.com/Bearded-Hen/Android-Bootstrap>



- **Ion**

make it easier to download files from the web.
<https://github.com/koush/ion>

- **ButterKnife**

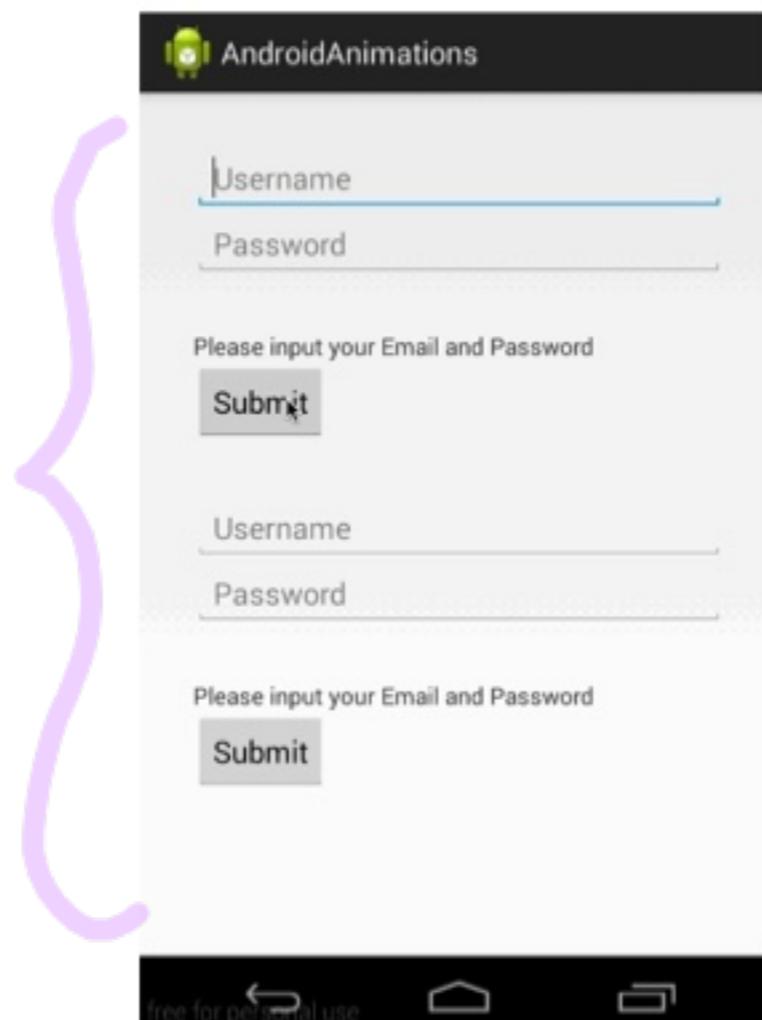
Popular library intended to simplify usage of Android widgets and events in Java code.
<http://jakewharton.github.io/butterknife/>



- Ambitious Android user named **daimajia** has created several libraries, including one to do animation effects on Views.

<https://github.com/daimajia/AndroidViewAnimations>

Animation Library



Graph View

Open-source graph plotting library for Android



- Adding GraphView to your project:

Add that line to your ***build.gradle*** file under your **app** directory into the dependencies block:

```
implementation 'com.jjoe64:graphview:4.2.2'
```

Graph View: Line graph

```
fun lineplot(){
    val series = LineGraphSeries(
        arrayOf(
            DataPoint(0.0, 1.0),
            DataPoint(1.0, 5.0),
            DataPoint(2.0, 3.0),
            DataPoint(3.0, 2.0),
            DataPoint(4.0, 6.0)
        )
    )
    mGraphX.addSeries(series)
}
```

Graph View: Bar graph

```
fun barplot(){
    val series = BarGraphSeries(
        arrayOf(
            DataPoint(0.0, -1.0),
            DataPoint(1.0, 5.0),
            DataPoint(2.0, 3.0),
            DataPoint(3.0, 2.0),
            DataPoint(4.0, 6.0)
        )
    )
    mGraphX.addSeries(series)

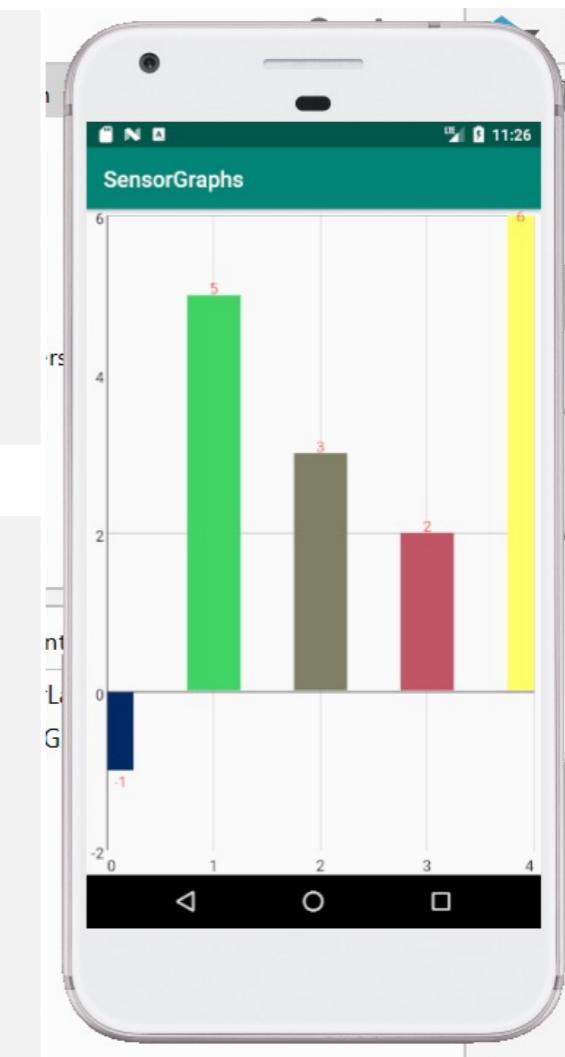
    // styling
    series.setValueDependentColor {data ->Color.rgb(data.x.toInt() * 255 / 4,
        Math.abs(data.y * 255 / 6).toInt(),100
    )
}
series.spacing = 50

// draw values on top
series.isDrawValuesOnTop = true
series.valuesOnTopColor = Color.RED
}
```

Graph View

```
fun barplot(){ //call this function anywhere, on the onCreate for example
    val series = BarGraphSeries(
        arrayOf(
            DataPoint(0.0, -1.0),
            DataPoint(1.0, 5.0),
            DataPoint(2.0, 3.0),
            DataPoint(3.0, 2.0),
            DataPoint(4.0, 6.0)      )
    )
    mGraphX.addSeries(series)

<?xml version="1.0" encoding="utf-8"?>
<LinearLayout xmlns:android="http://schemas.android.com/apk/res/android"
    xmlns:tools="http://schemas.android.com/tools"
    xmlns:app="http://schemas.android.com/apk/res-auto"
    android:layout_width="match_parent"
    android:layout_height="match_parent"
    android:orientation="vertical"
    tools:context=".MainActivity">
    <com.jjoe64.graphview.GraphView
        android:layout_weight="1"
        android:layout_width="match_parent"
        android:layout_height="wrap_content"
        android:id="@+id/mGraphX"/>
</LinearLayout>
```



Graph View: Real Time Graph

```
override fun onCreate(savedInstanceState: Bundle?) {
    super.onCreate(savedInstanceState)

    mSeriesXaccel = LineGraphSeries()
    mSeriesYaccel = LineGraphSeries()
    initGraphRT(mGraphX,mSeriesXaccel!!)
    initGraphRT(mGraphY,mSeriesYaccel!!)
}
```

```
override fun onSensorChanged(event: SensorEvent) {

    if (event.sensor.type != Sensor.TYPE_ACCELEROMETER)
        return
    linear_acceleration[0] = event.values[0]
    linear_acceleration[1] = event.values[1]
    linear_acceleration[2] = event.values[2]

    val xval = System.currentTimeMillis()/1000.toDouble()//graphLastXValue += 0.1
    mSeriesXaccel!!.appendData(DataPoint(xval, linear_acceleration[0].toDouble()),
    true, 50)
    mSeriesYaccel!!.appendData(DataPoint(xval, linear_acceleration[1].toDouble()),
    true, 50)
}
```

Face Detection with MLKit

EE P 523, Lecture 5



Mobile and web application development platform

In this course we will learn about and use

- Remote databases (Lecture 6-7)
- Machine Learning tools -> **Face Detection**



Face Detection: *MLkit*



Use Firebase for Android

Add Firebase to your Android App:

1. Create a project in your Firebase console
2. Add an Android app
 - a) from Android Studio OR
 - b) from firebase cosole

<https://firebase.google.com/docs/android/setup#assistant>

Add Firebase to Android App

1. Go to Firebase online and create a new account
2. Add new project
3. Create an Android App where you are going to use Firebase
4. Add Firebase to your Android App
5. Connect your App to Firebase
In Android IDE, go to [Tools -> Firebase ->](#)
6. Follow the steps

Add Project in Firebase console

The image shows two screenshots of the Firebase console. On the left is the main dashboard with a 'Recent projects' section and a prominent 'Add project' button. On the right is a detailed 'Add a project' dialog.

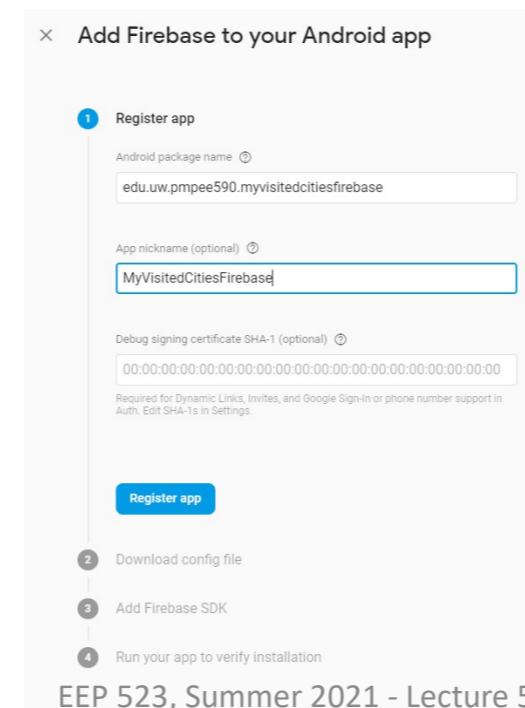
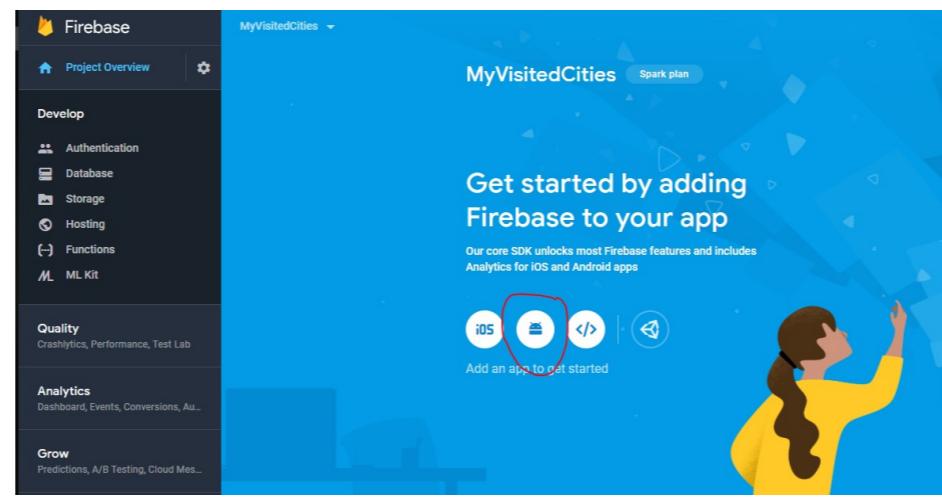
Add a project Dialog:

- Project name:** MyVisitedCities
- Tip:** Projects span apps across platforms
- Project ID:** my-visited-cities
- Locations:** United States (Analytics), nam5 (us-central) (Cloud Firestore)
- Analytics sharing settings:** A list of checked items:
 - Share your Analytics data with all Firebase features
 - Share your Analytics data with Google to improve Google Products and Services
 - Share your Analytics data with Google to enable technical support
 - Share your Analytics data with Google to enable Benchmarking
 - Share your Analytics data with Google Account Specialists
- Acceptance terms:** I accept the [controller-controller terms](#). This is required when sharing Analytics data to improve Google Products and Services. [Learn more](#)
- Buttons:** Cancel, Create project

7/29/21

EEP 523, Summer 2021 - Lecture 5

Add App to Firebase project



7/29/21

EEP 523, Summer 2021 - Lecture 5

Add App to Firebase project

X Add Firebase to your Android app

1 Register app
Android package name: edu.uw.pmppe590.myvisitedcitiesfirebase, App nickname: MyVisitedCitiesFirebase

2 Download config file

[Download google-services.json](#)

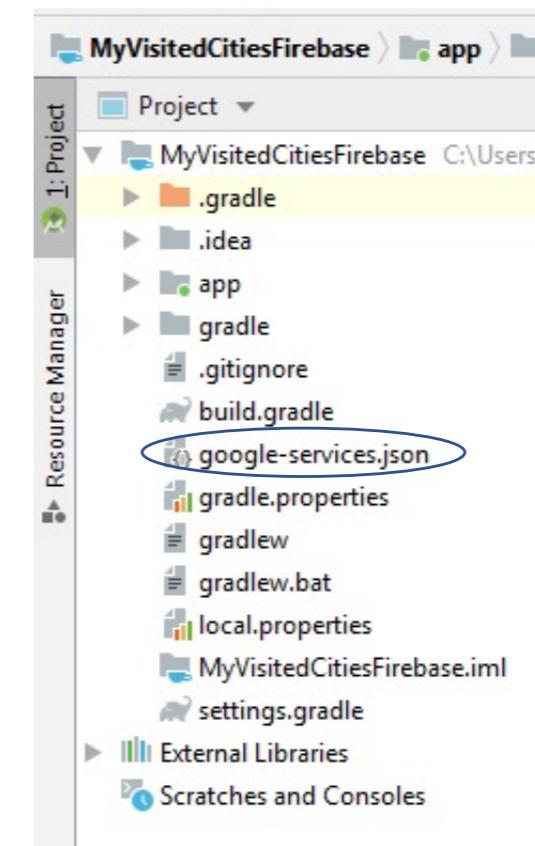
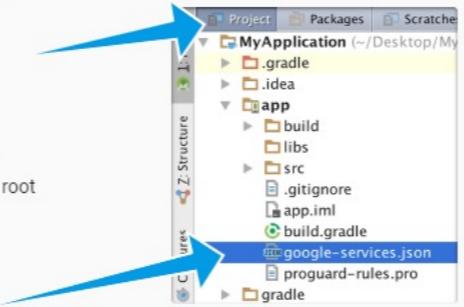
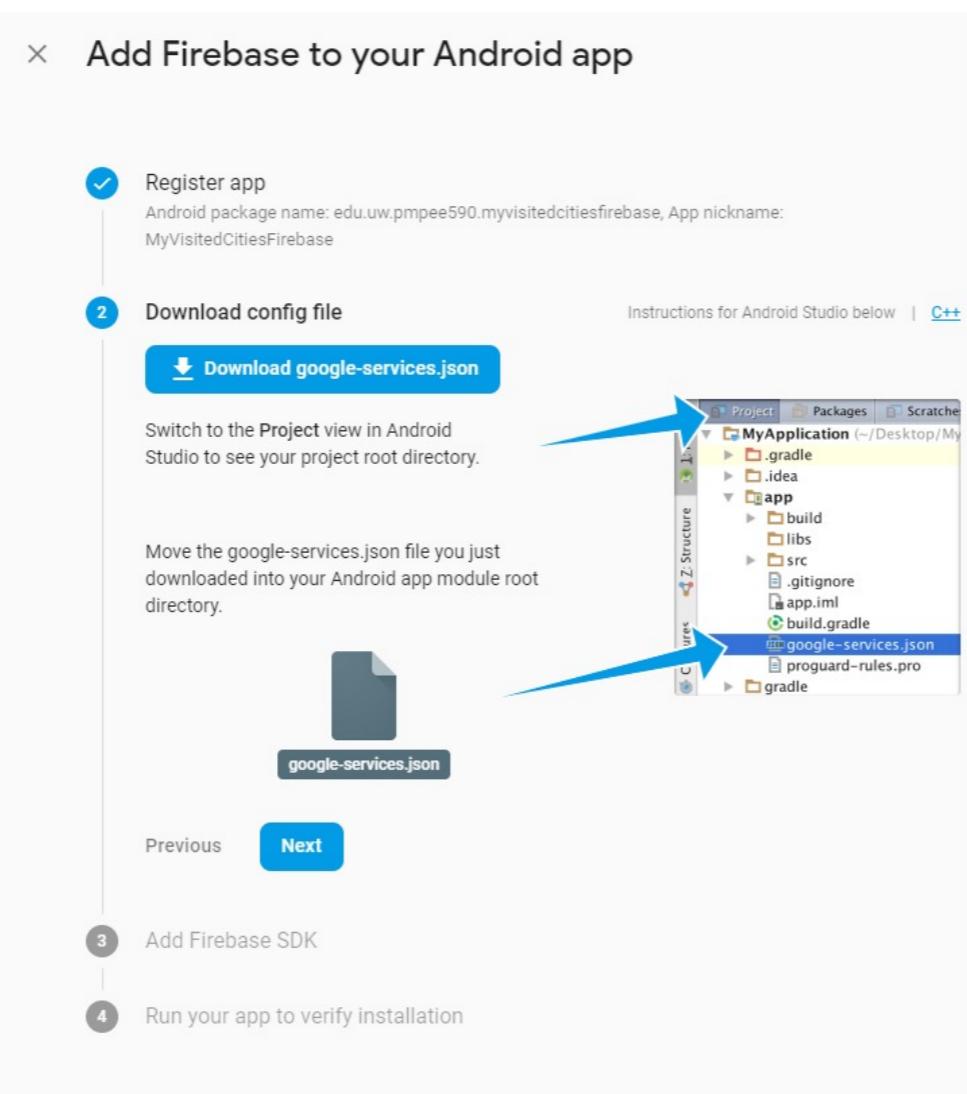
Switch to the Project view in Android Studio to see your project root directory.

Move the google-services.json file you just downloaded into your Android app module root directory.

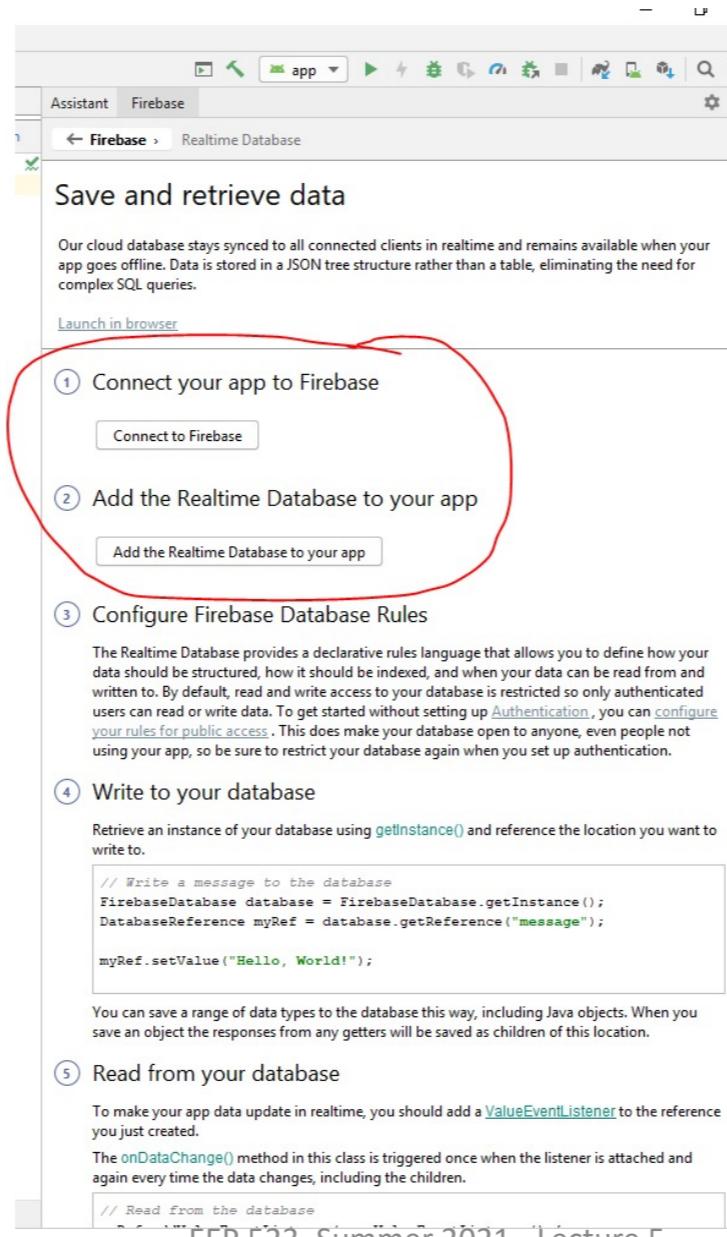
Next

3 Add Firebase SDK

4 Run your app to verify installation



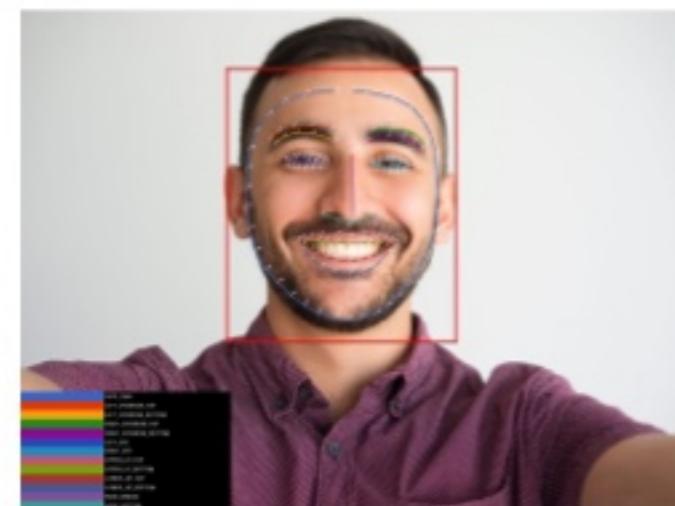
From AndroidStudio Assistant



Face Detection: *Contour points*

Each feature contour that ML Kit detects is represented by a fixed number of points:

<u>Face oval</u>	36 points	<u>Upper lip (top)</u>	11 points
<u>Left eyebrow (top)</u>	5 points	<u>Upper lip (bottom)</u>	9 points
<u>Left eyebrow (bottom)</u>	5 points	<u>Lower lip (top)</u>	9 points
<u>Right eyebrow (top)</u>	5 points	<u>Lower lip (bottom)</u>	9 points
<u>Right eyebrow (bottom)</u>	5 points	<u>Nose bridge</u>	2 points
<u>Left eye</u>	16 points	<u>Nose bottom</u>	3 points
<u>Right eye</u>	16 points		
<u>Left cheek (center)</u>	1 point		
<u>Right cheek (center)</u>	1 points		

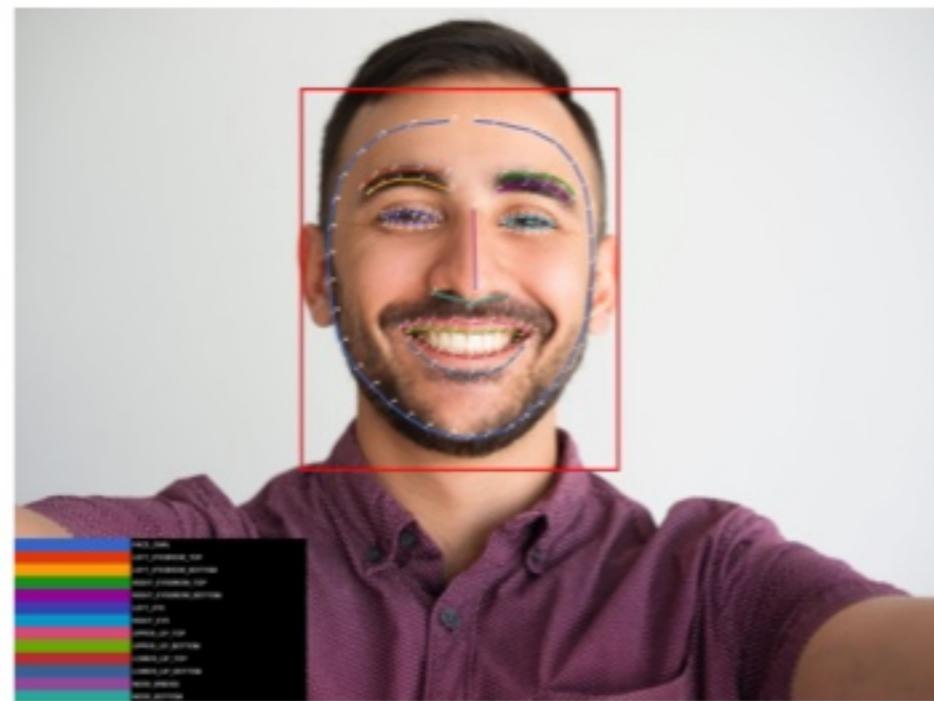


Indexes of feature contours	
0-35	Face oval
36-40	Left eyebrow (top)
41-45	Left eyebrow (bottom)
46-50	Right eyebrow (top)
51-55	Right eyebrow (bottom)
56-71	Left eye
72-87	Right eye
88-96	Upper lip (bottom)
97-105	Lower lip (top)
106-116	Upper lip (top)
117-125	Lower lip (bottom)
126, 127	Nose bridge
128-130	Nose bottom (note that the center point is at index 128)
131	Left cheek (center)
132	Right cheek (center)

Face Detection: *Landmarks*

A landmark is a point of interest within a face.

*right eye,
left eye,
nose base,
left cheek,
right cheek,
left mouth,
right mouth,
bottom mouth*



Face Detection: *Classification*

- Classification determines whether a certain facial characteristic is present.
- ML Kit currently supports two classifications: **eyes open** and **smiling**.
- Classification is expressed as a certainty value, indicating the **confidence** that the facial characteristic is present.
- For example, a value of 0.7 or more for the smiling classification indicates that it is likely that a person is smiling.
- Both of these classifications rely upon **landmark detection**.
- Also note that "eyes open" and "smiling" classification only works for **frontal faces**,

Face Detection: *Firebase MLkit*

Setting up your project to use the face detector of MLkit

1. If you haven't already, [add Firebase to your Android project](#).
2. In your project-level build.gradle file, make sure to include Google's Maven repository in both your buildscript and allprojects sections.
3. Add the dependencies for the ML Kit Android libraries to your module (app-level) Gradle file ([app/build.gradle](#))

```
apply plugin: 'com.google.gms.google-services'

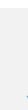
dependencies {
    // ...

    implementation 'com.google.firebase:firebase-ml-vision:24.0.2'
    // If you want to detect face contours (landmark detection and classification
    // don't require this additional model):
    implementation 'com.google.firebase:firebase-ml-vision-face-model:20.0.0'
}
```

Face Detection: *Firebase MLkit*

```
private fun faceDetection() {  
  
    val image = FirebaseVisionImage.fromBitmap(mSelectedImage!!)  
    val options = FirebaseVisionFaceDetectorOptions.Builder()  
        .setPerformanceMode(FirebaseVisionFaceDetectorOptions.FAST)  
        .setLandmarkMode(FirebaseVisionFaceDetectorOptions.ALL_LANDMARKS)  
        .setContourMode(FirebaseVisionFaceDetectorOptions.ALL_CONTOURS)  
        .setClassificationMode(FirebaseVisionFaceDetectorOptions.ALL_CLASSIFICATIONS)  
        .build()  
    val detector = FirebaseVision.getInstance().getVisionFaceDetector(options)  
    detector.detectInImage(image)  
        .addOnSuccessListener { faces ->  
            processFaces(faces)  
        }  
        .addOnFailureListener { e -> // Task failed with an exception  
            e.printStackTrace()  
        }  
}
```

Bitmap image



Face Detection: *Firebase MLkit*

Landmarks Example: get the pixel positions (x,y) of the face center and eyes

```
private fun processFaces(faces: List<FirebaseVisionFace>) {
    // Task completed successfully
    if (faces.size == 0) {
        return
    }

    val face = faces[0]
    val center_x = (face.boundingBox.centerX().toInt())
    val center_y = (face.boundingBox.centerY().toInt())

    var leftEyePosX = 0
    var leftEyePosY = 0
    val leftEye = face.getLandmark(FirebaseVisionFaceLandmark.LEFT_EYE)
    leftEye?.let {
        leftEyePosX = leftEye.position.x.toInt()
        leftEyePosY = leftEye.position.y.toInt()
    }
}
```

Face Detection: *Firebase MLkit*

Contour Example: get the pixel positions of the contour points

```
val contour = face.getContour(FirebaseVisionFaceContour.ALL_POINTS)  
for (point in contour.points) {  
    val px = point.x  
    val py = point.y }
```



int	ALL_POINTS	All points of a face contour.
int	FACE	The outline of the subject's face.
int	LEFT_EYE	The outline of the subject's left eye cavity.
int	LEFT_EYEBROW_BOTTOM	The bottom outline of the subject's left eyebrow.
int	LEFT_EYEBROW_TOP	The top outline of the subject's left eyebrow.
int	LOWER_LIP_BOTTOM	The bottom outline of the subject's lower lip.
int	LOWER_LIP_TOP	The top outline of the subject's lower lip.
int	NOSE_BOTTOM	The outline of the subject's nose bridge.
int	NOSE_BRIDGE	The outline of the subject's nose bridge.
int	RIGHT_EYE	The outline of the subject's right eye cavity.
int	RIGHT_EYEBROW_BOTTOM	The bottom outline of the subject's right eyebrow.
int	RIGHT_EYEBROW_TOP	The top outline of the subject's right eyebrow.
int	UPPER_LIP_BOTTOM	The bottom outline of the subject's upper lip.
int	UPPER_LIP_TOP	The top outline of the subject's upper lip.

Face Detection: *Firebase MLkit*

Classification Example: get the probability of smile and eyes open

```
if (face.smilingProbability != FirebaseVisionFace.UNCOMPUTED_PROBABILITY) {  
    val smileProb = face.smilingProbability  
}  
if (face.rightEyeOpenProbability != FirebaseVisionFace.UNCOMPUTED_PROBABILITY) {  
    val rightEyeOpenProb = face.rightEyeOpenProbability  
}
```

Speech to Text

EE P 523, Lecture 5

Speech-to-Text

User talks; Android records, turns into a String.

- Similar to the camera, Android has a built-in activity for capturing speech into text.
- You can call it using an Intent and wait for the result.

```
fun record(){
    val intent = Intent(RecognizerIntent.ACTION_RECOGNIZE_SPEECH )
    intent.putExtra( RecognizerIntent.EXTRA_LANGUAGE, Locale.getDefault() )
    intent.putExtra(RecognizerIntent.EXTRA_LANGUAGE_MODEL,RecognizerIntent.EXTRA_LANGUAGE_MODEL_FREE_FORM)
    // prompt text is shown on screen to tell user what to say
    intent.putExtra(RecognizerIntent.EXTRA_PROMPT, "text")
    startActivityForResult(intent, KEY)
}
```

Receiving the Speech-to-text()

When the speech-to-text activity comes back, its Intent gives you all text spoken by the user in an ArrayList.

- Usually the first element (index 0) contains the string you want.

```
override fun onActivityResult(requestCode: Int, resultCode: Int, intent: Intent?) {  
    super.onActivityResult(requestCode, resultCode, intent)  
    val list = intent!!.getStringArrayListExtra(RecognizerIntent.EXTRA_RESULTS )  
    val spokenText = list[0]  
    val duration = Toast.LENGTH_SHORT  
  
    val toast = Toast.makeText(applicationContext, spokenText, duration)  
    toast.show()  
    // ...  
}
```

Robust Speech-to-text()

- Some devices do not have speech-to-text capability.
 - In these cases, it will throw an exception when you try to use it.
 - To handle such situations, you can **try/catch** the exception.

```
fun record(){  
    val intent = Intent(RecognizerIntent.ACTION_RECOGNIZE_SPEECH)  
    intent.putExtra( RecognizerIntent.EXTRA_LANGUAGE, Locale.getDefault() )  
    intent.putExtra(RecognizerIntent.EXTRA_LANGUAGE_MODEL,RecognizerIntent.LANGUAGE_MODEL_FREE_FORM)  
    // prompt text is shown on screen to tell user what to say  
    intent.putExtra(RecognizerIntent.EXTRA_PROMPT, "text")
```

```
    try {  
        startActivityForResult(intent, KEY) } ↗  
    catch (anfe: ActivityNotFoundException) {  
        //code to handle exception  
    } }
```



Text to Speech

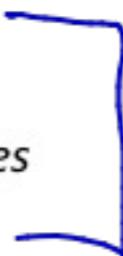
EE P 523, Lecture 5

Text-To-Speech

text-to-speech:

Allows Android device to speak an audible message based on a text string.

- Not installed by default on some devices.
To install, click
 - Settings → Language and Input → Text to speech output
 - Google textto-speech engine "settings" icon → Install voice data → Languages
- In general, text-to-speech on Android is simple:
 1. create a TextToSpeech object
 2. call its speak method



Text-To-Speech Class

TextToSpeech constructor and methods:

- `TextToSpeech(activity, Listener)` - constructor
 - `getVoice`, `getVoices`, `setVoice` - change speaking voice
 - `getLanguage`, `setLanguage` - sets language used
 - `getPitch`, `setPitch` - sets vocal tone used
 - `isSpeaking` - returns true if speaking
 - `shutdown` - kills TTS engine
- `speak(text, mode, params)` - speaks given text aloud
- `stop` - halts any speech
 - `synthesizeToFile(text, params, filename)` - speaks to file

The speak() method

```
speak(CharSequence text, int queueMode, Bundle params, String utteranceId)
```

The speak method accepts 4 parameters *:

- the text to speak aloud, as a String.
- the mode to use for speaking, one of:
 - TextToSpeech.`QUEUE_ADD` : Speak after any other text is done.
 - TextToSpeech.`QUEUE_FLUSH` : Stop any other text and speak immediately.
- a Map of parameters (we don't need any, so we can pass null).
- an ID of the utterance

* (Android 5 introduces a different version of speak(), but we'll ignore it.)

```
private fun speakOut() {  
    val myText1 = "Did you sleep well?"  
    val myText2 = "I hope so, because it's time to wake up."  
    tts.speak(myText1, TextToSpeech.QUEUE_FLUSH, null, null)  
    tts.speak(myText2, TextToSpeech.QUEUE_ADD, null, null)
```

Text-to-Speech

```
class MainActivity : AppCompatActivity(), TextToSpeech.OnInitListener {
    private lateinit var tts: TextToSpeech;

    override fun onCreate(savedInstanceState: Bundle?) {
        super.onCreate(savedInstanceState)
        setContentView(R.layout.activity_main)

        tts = TextToSpeech(this, this)
    }
}
```

```
override fun onInit(status: Int) {
    if (status == TextToSpeech.SUCCESS) {
        val result = tts!!.setLanguage(Locale.US) // set US English as Language for tts
        if (result == TextToSpeech.LANG_MISSING_DATA || result == TextToSpeech.LANG_NOT_SUPPORTED)
            Log.e("TTS", "The Language specified is not supported!")
    } else {
        Log.e("TTS", "Initialization Failed!")
    }
}
```

```
private fun speakOut() {
    val myText1 = "Did you sleep well?"
    val myText2 = "I hope so, because it's time to wake up."
    tts.speak(myText1, TextToSpeech.QUEUE_FLUSH, null, null)
    tts.speak(myText2, TextToSpeech.QUEUE_ADD, null, null)
}
```

Tiny Deep Learning: Practical Introduction

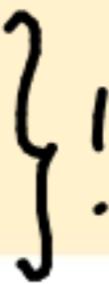
EE P 523, Lecture 5

Motivation and Goals for This Topic

- Detect your own custom gesture (not just known patterns)
- Accurate and ~fast data classification and prediction

There are many books and courses that explain the science behind deep learning, so we won't be doing that here
(encourage you to explore!)

Remember:
you don't need all the theory to
start building useful things



What is AI/ML/DL

ARTIFICIAL INTELLIGENCE

Technique that
enables computers
to mimic human
behavior



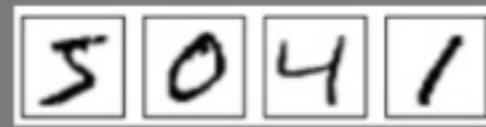
MACHINE LEARNING

Ability to *learn*
without explicitly
being programmed



DEEP LEARNING

Extract patterns from data
using Neural Networks

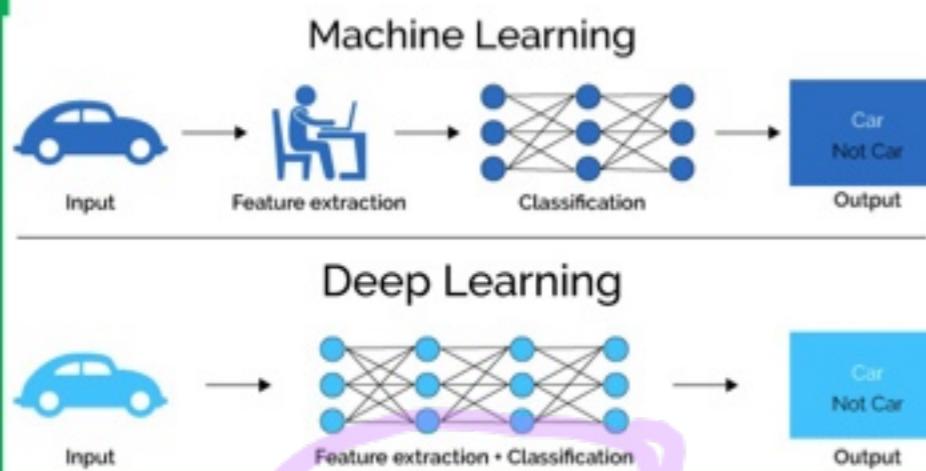


↑
MNIST

Deep Learning

WHY ?

- What is it: pattern from
Extract useful patterns from data
- How:
neural network + optimization
- How in practice:
Python + TensorFlow
- Why now:
data, hardware, community,
tools, investment



WHY NOT ?



Unintended
consequences



Three Classes of Learning Problems

Supervised

Data (x, y)

x : data

y : labels

Goal learn function to map

$$x \rightarrow y$$

Example:

This is a digit 3



Unsupervised

Data x (no labels!)

Goal: learn underlying structure

Example:

This thing is like the other thing



Reinforcement

Data state-action pairs

Goal: maximize future rewards

Example:

Select this number because it will make you go forward

Getting Up to Speed on Machine Learning

Creating a machine learning program is different from the usual process of writing code.

- In a traditional piece of software, a programmer designs an algorithm that takes an input, applies various rules, and returns an output.
- The algorithm's internal operations are planned out by the programmer and implemented explicitly through lines of code
- To create a machine learning program, a programmer feeds data into a special kind of algorithm and lets the algorithm discover the rules.
- This means that, as programmers, we can create programs that make predictions based on complex data without having to understand of all the complexity ourselves

Deep Learning Workflow in 7 steps

1. Decide on a goal
2. Collect a dataset
-  3. Design a model architecture
4. Train the model
5. Convert the model
6. Run inference
7. Evaluate and troubleshoot

Images are Numbers



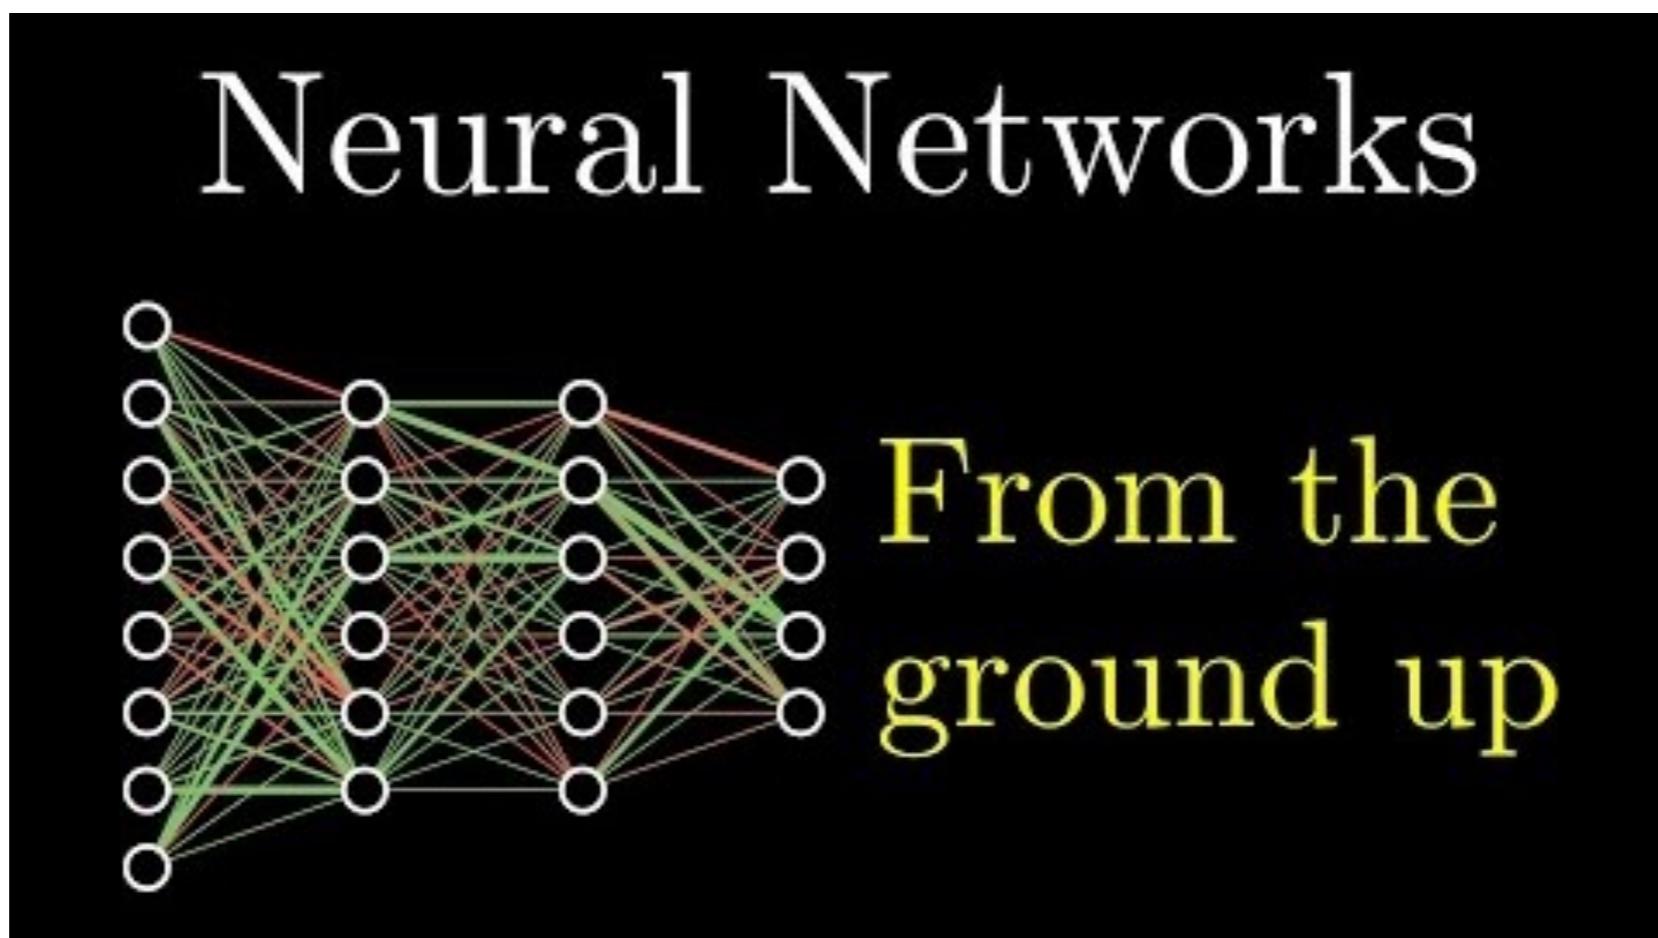
167	163	174	168	160	162	129	161	172	143	156	156
166	162	163	74	75	62	33	17	110	210	180	154
180	180	80	14	34	6	10	33	48	106	169	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	261	237	239	239	238	227	87	71	201
172	106	207	239	239	214	220	239	228	98	74	206
188	88	179	209	185	215	211	168	139	75	20	169
189	97	166	84	10	168	134	11	31	62	22	148
199	148	191	189	188	227	178	143	162	106	36	190
206	174	166	282	236	231	149	178	228	43	95	234
190	216	136	149	236	187	86	150	78	38	218	241
190	234	147	196	227	210	137	102	36	105	258	234
190	214	173	66	109	143	96	50	2	100	249	219
187	196	236	75	1	81	47	0	6	217	256	211
183	202	237	145	0	0	12	108	200	138	243	236
196	206	123	207	177	121	120	209	179	11	96	218

An image is just a matrix of numbers [0,255]!
i.e., 1080 × 1080 × 3 for an RGB image

What the computer sees

167	163	174	168	160	162	129	161	172	161	156	156
166	162	163	74	75	62	33	17	110	210	180	154
180	180	80	14	34	6	10	33	48	106	169	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	261	237	239	239	238	227	87	71	201
172	106	207	239	239	214	220	239	228	98	74	206
188	88	179	209	185	215	211	168	139	75	20	169
189	97	166	84	10	168	134	11	31	62	22	148
199	168	191	189	188	227	178	143	162	106	36	190
206	174	166	282	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	78	38	218	241
190	234	147	196	227	210	137	102	36	105	258	234
190	214	173	66	109	143	96	50	2	100	249	219
187	196	236	75	1	81	47	0	6	217	256	211
183	202	237	145	0	0	12	108	200	138	243	236
196	206	123	207	177	121	120	209	179	11	96	218

1. Decide on a goal



1. Decide on the goal

- We can express this as a classification problem.
- Classification is a machine learning task that takes a set of input data and returns the probability that this data fits each of a set of known *classes*.

2. Collect Dataset

1. Select data

It's really important that the data you choose will also be available when you want to make predictions

2. Collect data

It's difficult to know exactly how much data is required to train an effective model. It depends on many factors, such as the complexity of the relationships between variables, the amount of noise, and the ease with which classes can be distinguished. However, there's a rule of thumb that is always true: the more data, the better!

3. Label data

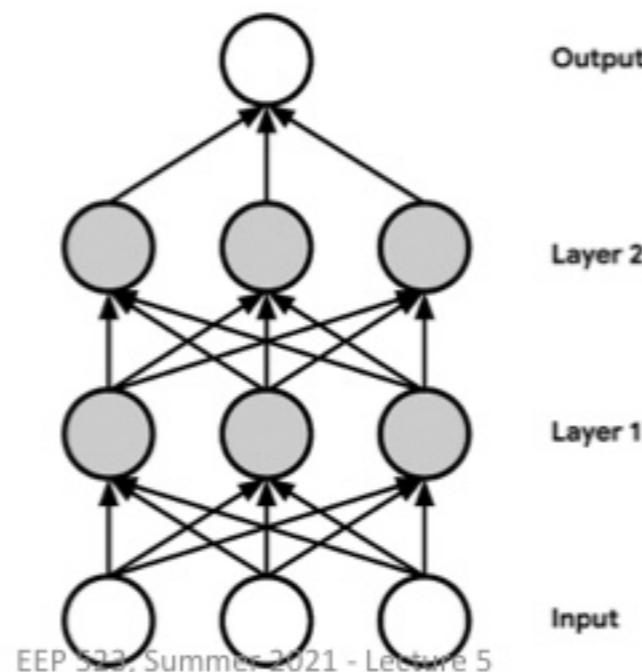
The process of associating data with classes is called labeling, and the “normal” and “abnormal” classes are our labels.

3. Design a Model Architecture

Model

- network of simulated neurons represented by arrays of numbers arranged in layers.
- These numbers are known as weights and biases, or collectively as the network's *parameters*.

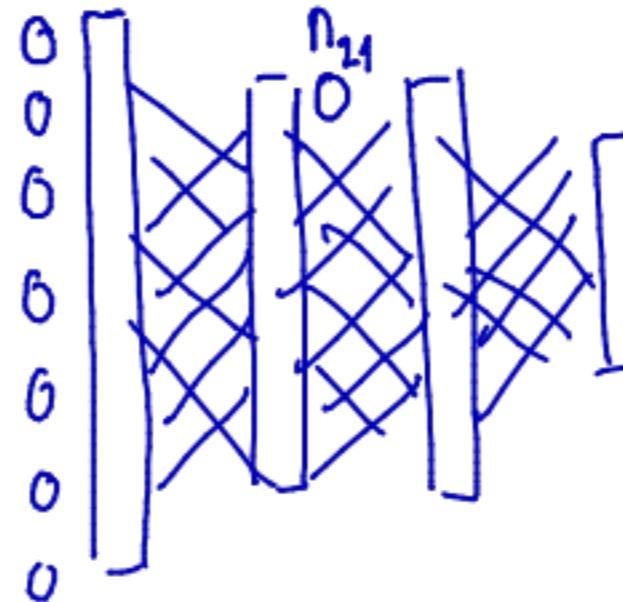
When data is fed into the network, it is transformed by successive mathematical operations that involve the weights and biases in each layer. The output of the model is the result of running the input through these operations.



Single Neuron (Perceptron)

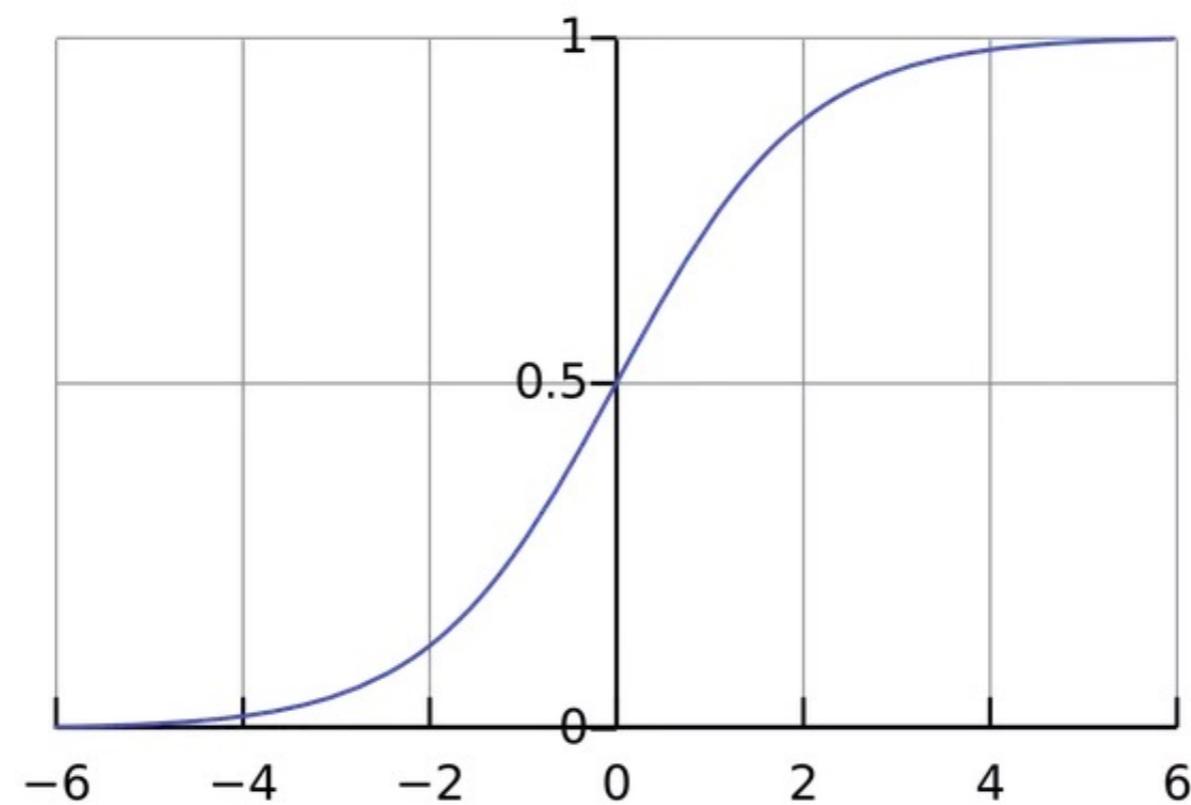
$$n_{21} = f(n_{1i} w)$$

Activation Function: ReLU

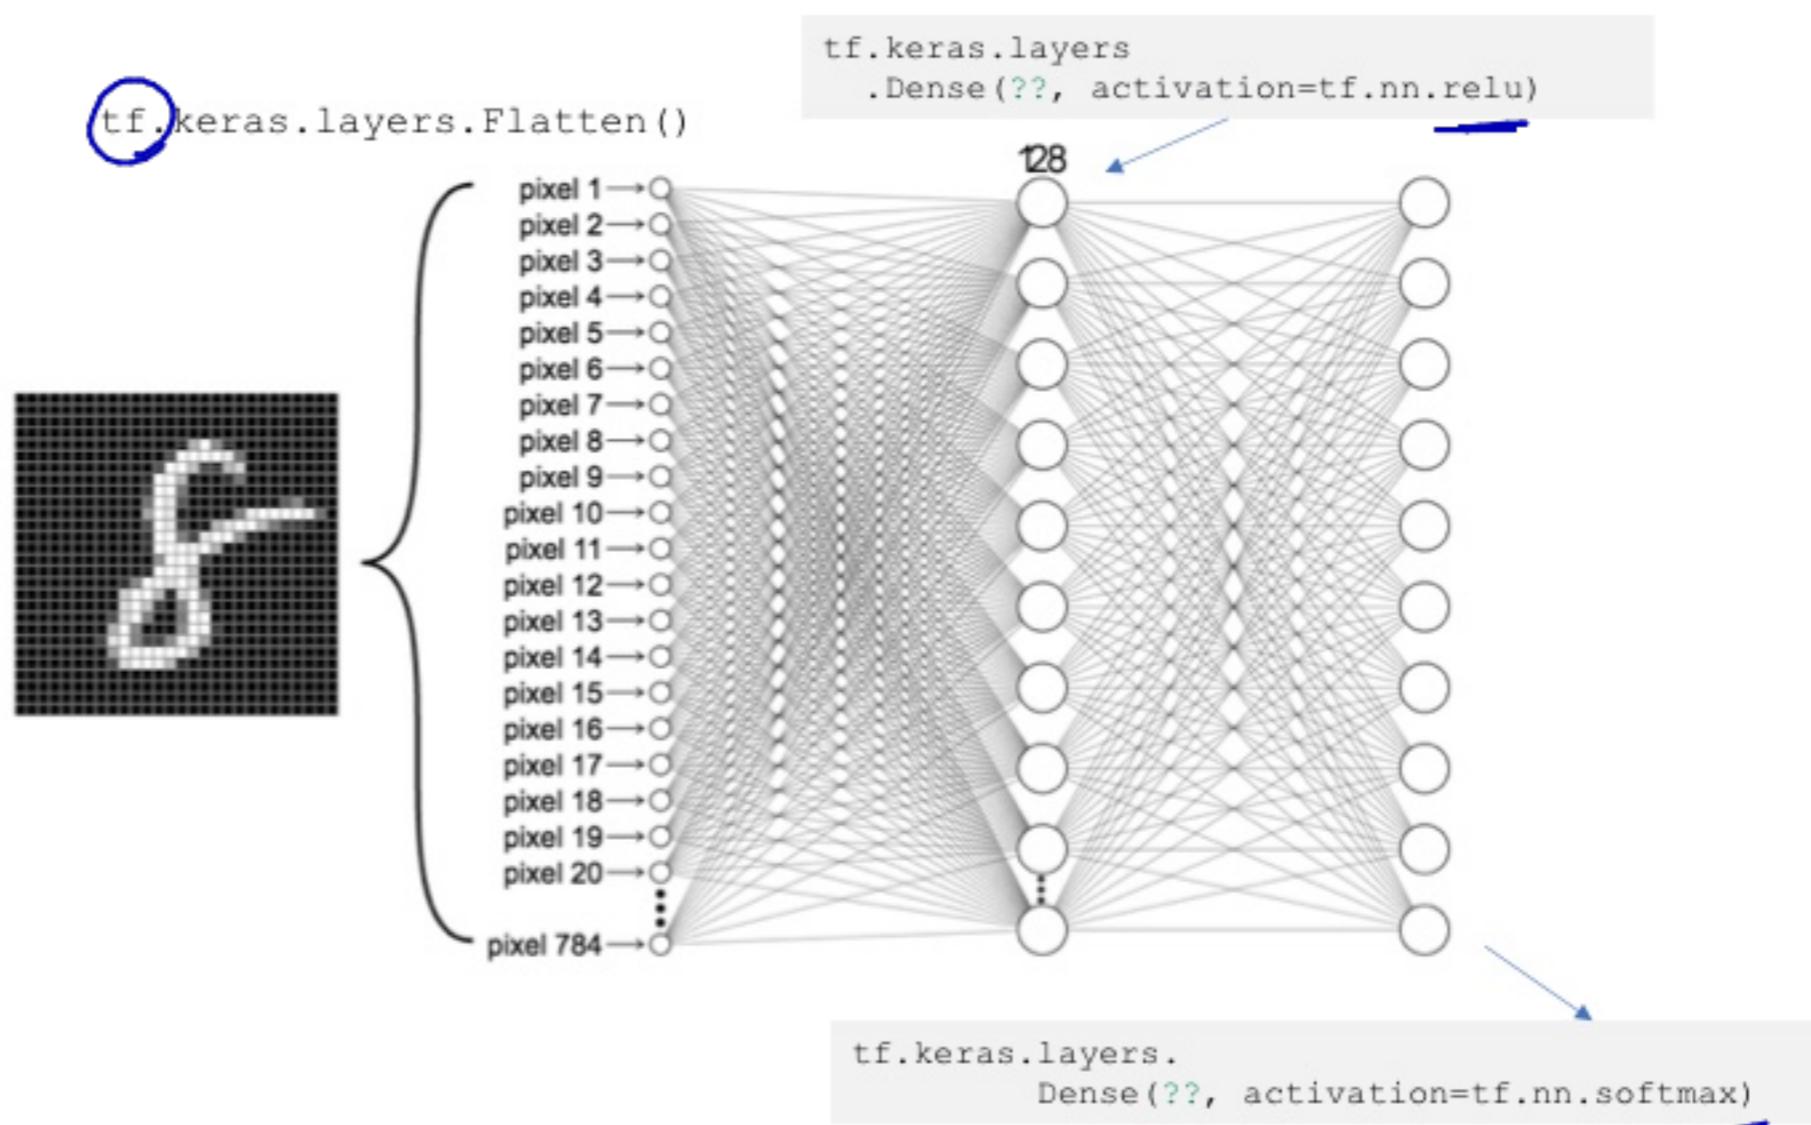


ReLU	Leaky ReLU	ELU
$g(z) = \max(0, z)$	$g(z) = \max(\epsilon z, z)$ with $\epsilon \ll 1$	$g(z) = \max(\alpha(e^z - 1), z)$ with $\alpha \ll 1$
• Non-linearity complexities biologically interpretable	• Addresses dying ReLU issue for negative values	• Differentiable everywhere

Activation Function: SoftMax



Our model



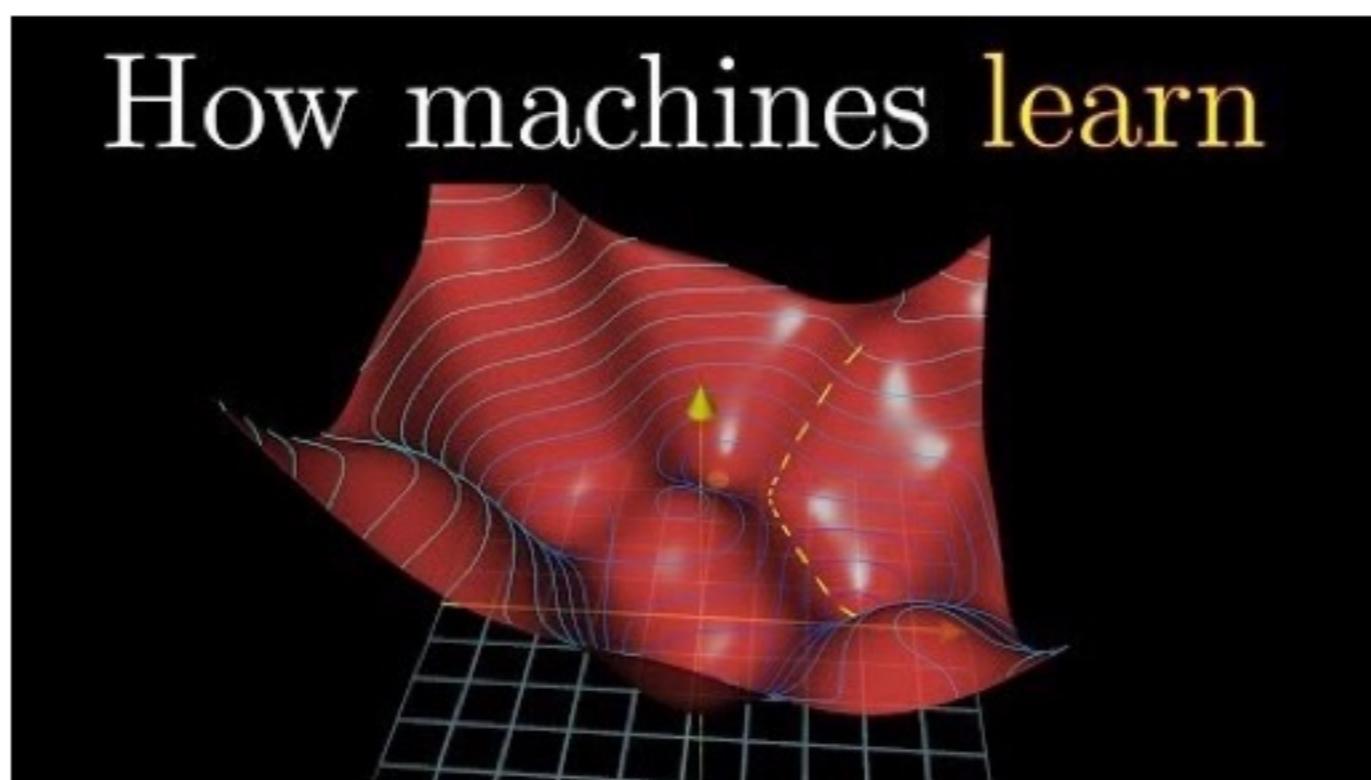
3. Design a Model Architecture

- For many common problems, you can find pretrained models available online for free.
- Deep learning models accept input and generate output in the form of tensors.
- For the purposes of this course, a tensor is essentially a list that can contain either numbers or other tensors; you can think of it as similar to an array.

4. Train the model

- The model's weights start out with random values, and biases typically start with a value of 0.
- During training, batches of data are fed into the model, and the model's output is compared with the desired output (which in our case is the correct label, "normal" or "abnormal").
- An algorithm called backpropagation adjusts the weights and biases incrementally so that over time, the output of the model gets closer to matching the desired value.
- Training, which is measured in epochs (meaning iterations), continues until we decide to stop.

4. How our Models Learn?

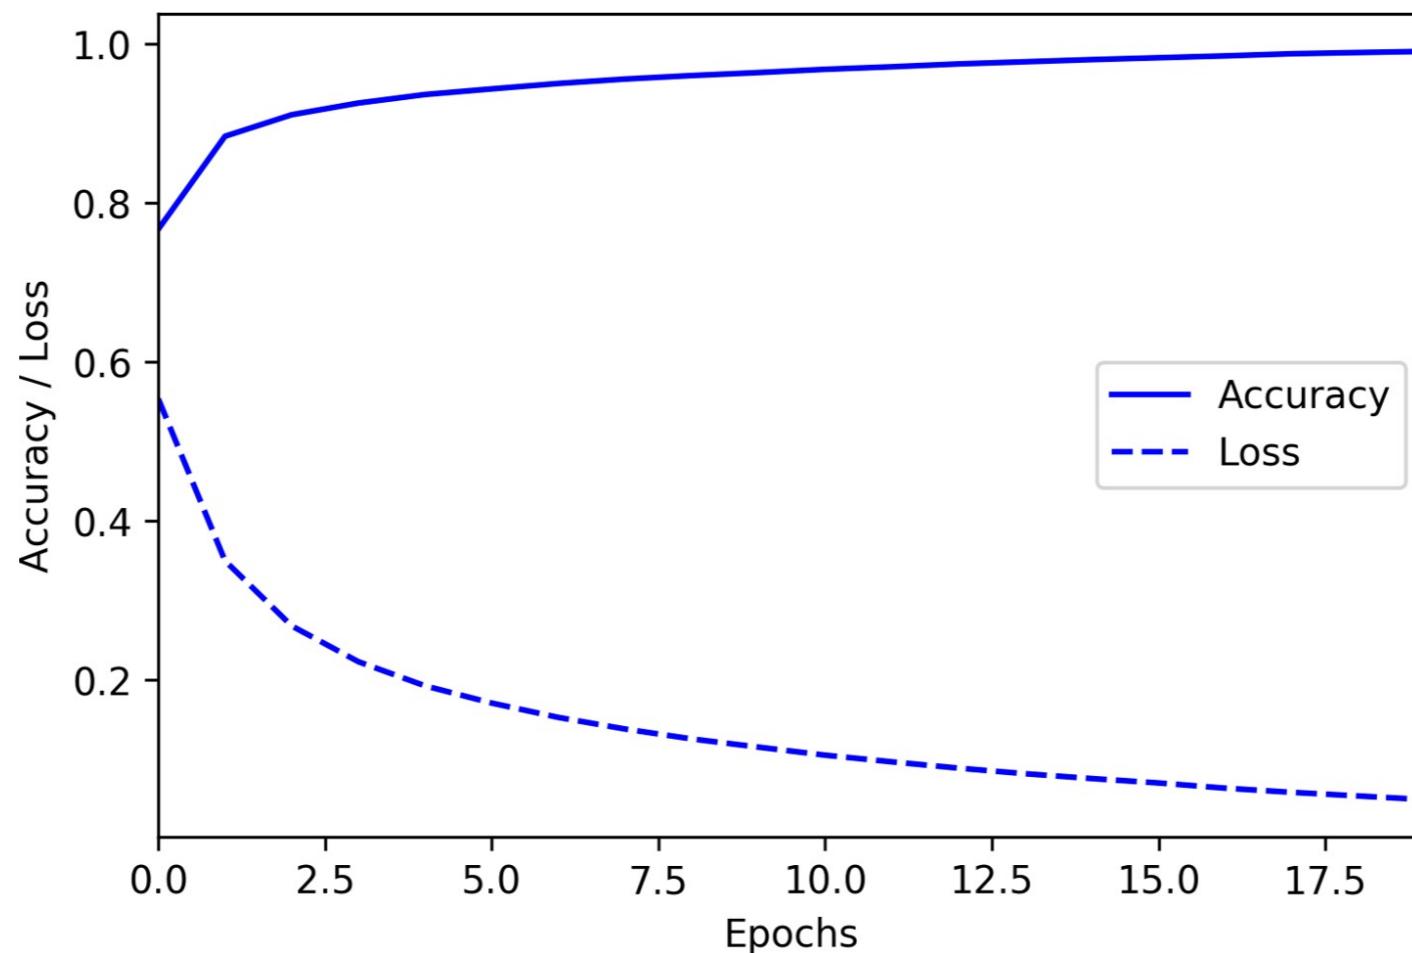


4. Train the Model: When to Stop?

- We generally stop training when a model's performance stops improving. It is said to have converged.
- Common performance metrics are *loss* and *accuracy*.
 - **Loss**: numerical estimate of how far the model is from producing the expected answers
 - **Accuracy**: percentage of the time that it chooses the correct prediction.
 - A perfect model would have a loss of 0.0 and an accuracy of 100%, but real models are rarely perfect.

4. Train the Model:

Loss and Accuracy



4. Train the Model: Hyperparameters

- To attempt to improve the model's performance, we can change our model architecture, and we can adjust various values used to set up the model and moderate the training process.
- These values are collectively known as *hyperparameters*, and they include variables such as the number of training **epochs** to run and the number of neurons in each layer.
- Each time we make a change, we can retrain the model, look at the metrics, and decide whether to optimize further.
- Hopefully, time and iterations will result in a model with acceptable accuracy!

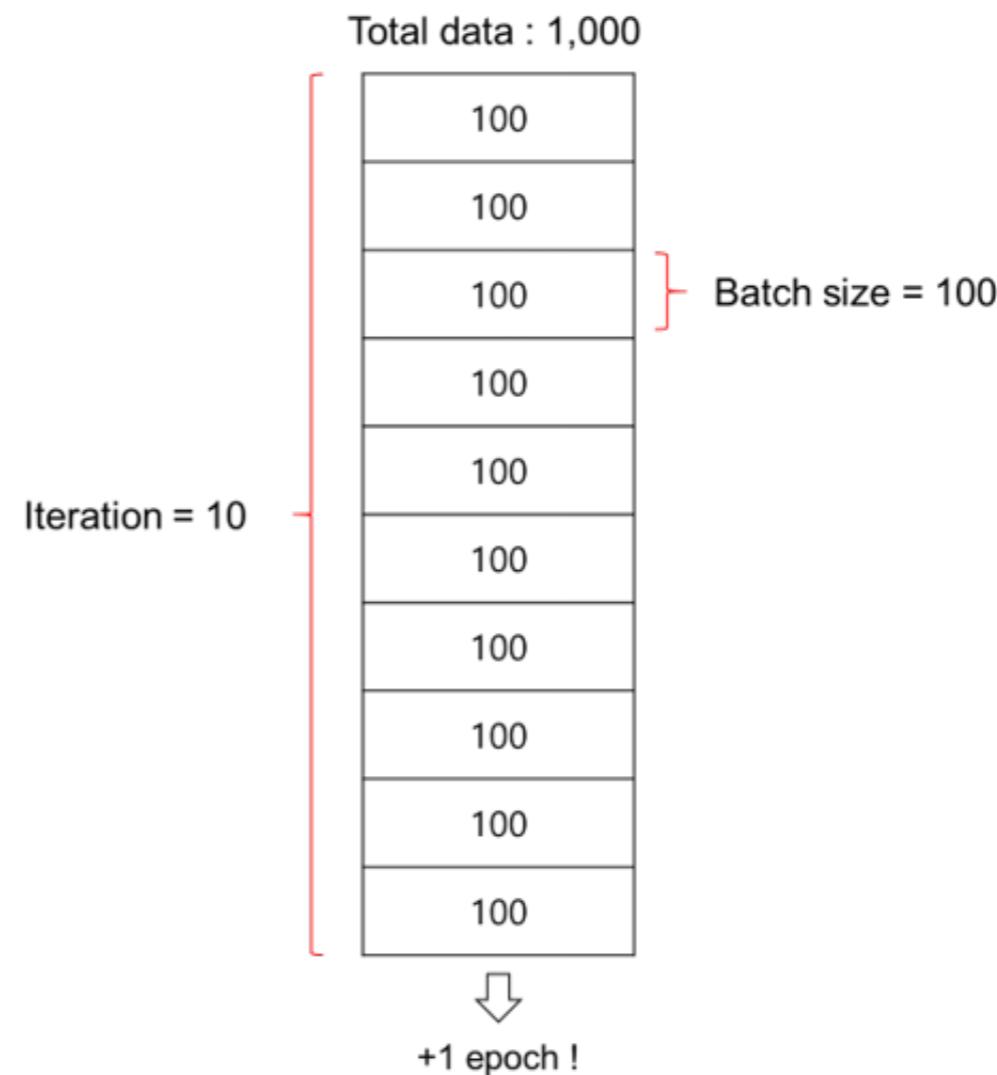
4. Train the Model: Hyperparameters

- **Batch size**

how many pieces of training data to feed into the network before measuring its accuracy and updating its weights and biases

- **Epochs**

how many times our entire training set will be run through the network during training. The more epochs, the more training will occur.



4

4. Train the Model: Why the Model Fails?

A neural network learns to *fit* its behavior to the patterns it recognizes in data

- ***Underfit***

model has not yet been able to learn a strong enough representation of these patterns to be able to make good predictions.

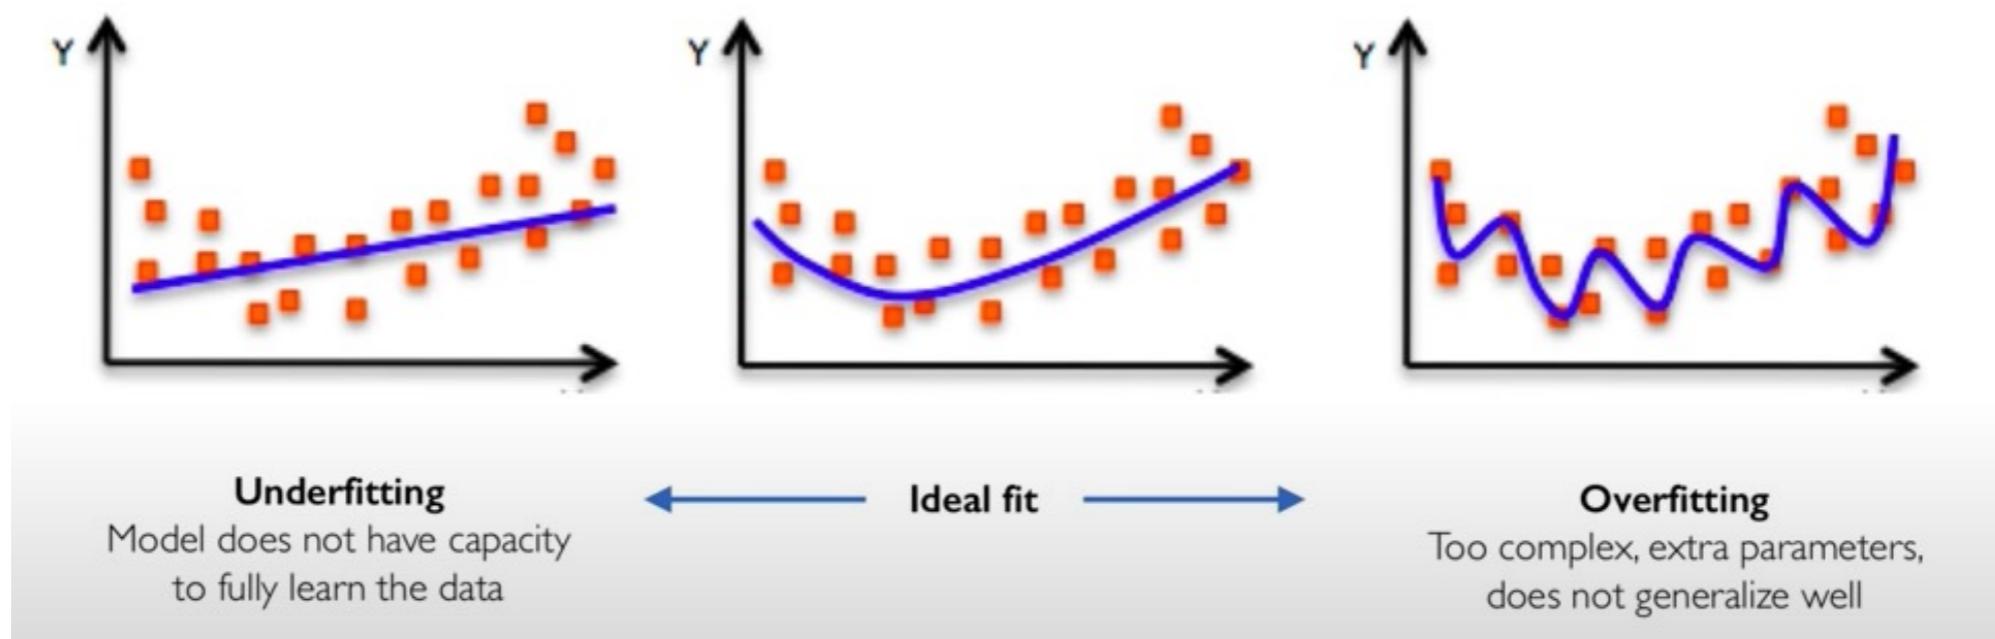
- ✗ The architecture is too small to capture the complexity of the system it is supposed to model
- ✗ The model has not been trained on enough data.

- ***Overfit***

it has learned its training data too well. The model is able to exactly predict the minutiae of its training data, but it is not able to generalize its learning to data it has not previously seen.

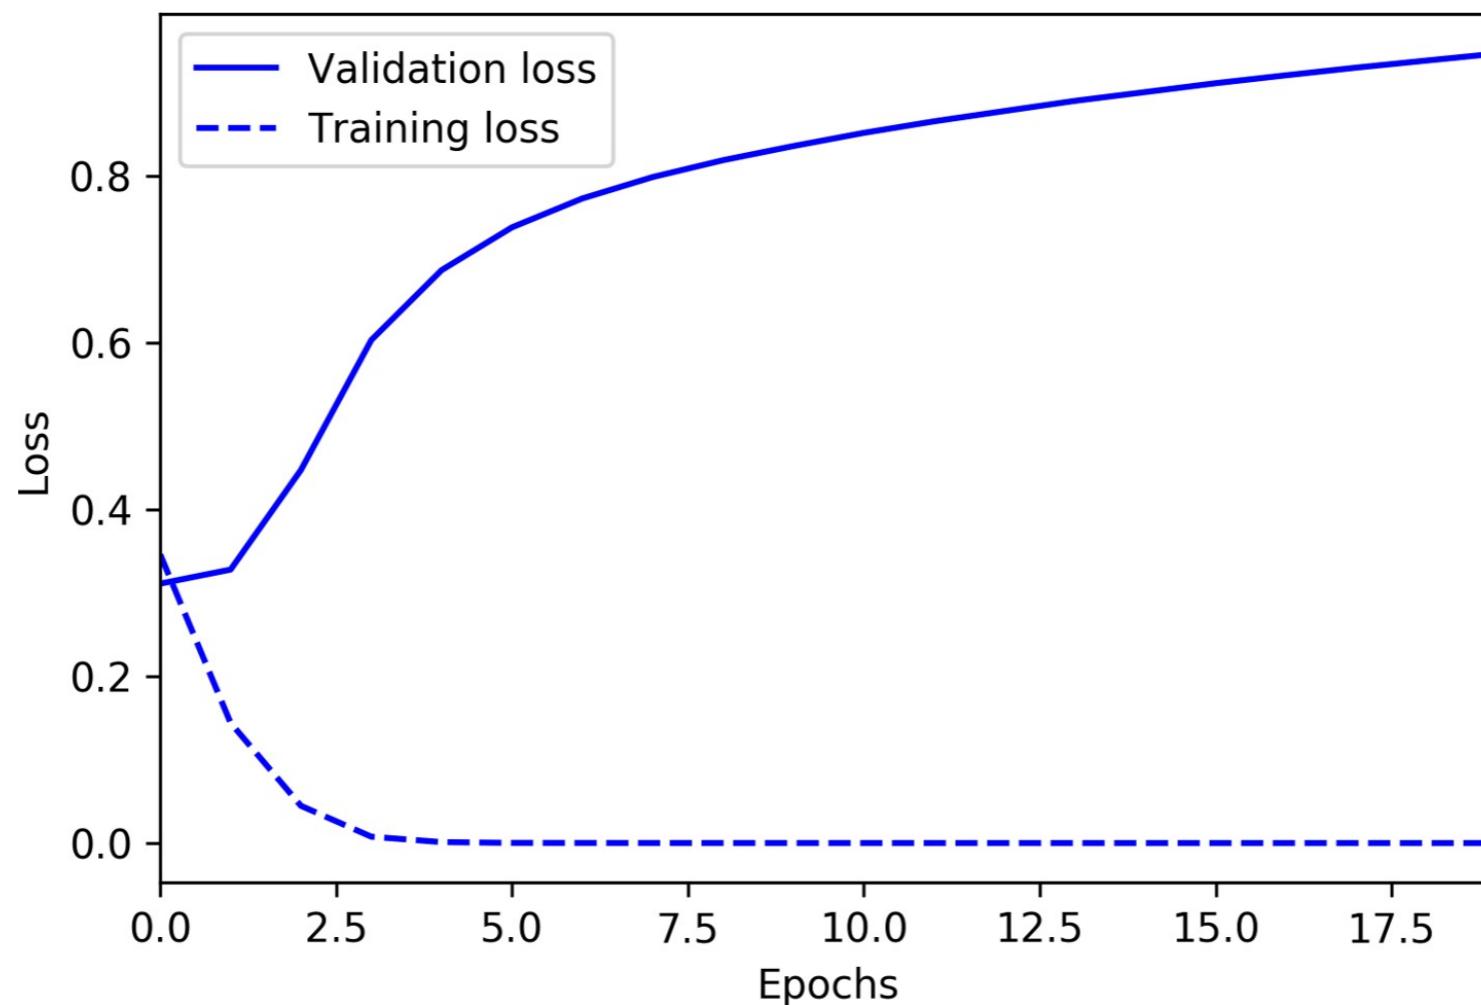
- ✗ the model has managed to entirely memorize the training data
- ✗ it has learned to rely on a shortcut present in the training data but not in the real world.

4. Train the model: Why the Model Fails?



4. Train the Model: Why the Model Fails?

Overfitting example



4. Train the model: Preventing the model from failing

Goal: make deep learning models **less likely to overfit their training data.**

Regularization

- generally involve constraining the model in some way in order to prevent it from perfectly memorizing the data that it's fed during training.
- Example: *dropout* -> randomly cutting the connections between neurons during training.

Data augmentation,

- artificially expand the size of a training dataset.
- creating multiple additional versions of every training input, each transformed in a way that preserves its meaning but varies its exact composition

4. Train the model: train-validate-test

- need to *validate* the model using new data that wasn't used in training. It's common to split a dataset into three parts—***training, validation, and test.***
- A typical split is **60% training data, 20% validation, and 20% test**. This splitting must be done so that each part contains the same distribution of information, and in a way that preserves the structure of the data.
 - 1) During training, the *training* dataset is used to train the model.
 - 2) Periodically, data from the *validation* dataset is fed through the model, and the loss is calculated. Because the model has not seen this data before, its loss score is a more reliable measure of how the model is performing.
 - 3) By comparing the training and validation loss (and accuracy, or whichever other metrics are available) over time, you can see whether the model is overfitting.

Deploying machine learning models on mobile and IoT devices

TensorFlow

- a set of tools for building, training, evaluating, and deploying machine learning models.
- Originally developed at Google - now an open source project
- most popular and widely used framework for machine learning.

TensorFlow Lite

- set of tools for deploying TensorFlow models to mobile and embedded devices

Keras

- TensorFlow's high-level API that makes it easy to build and train deep learning networks. We'll also use.



5. Convert the model: Lite model

TensorFlow model:

- set of instructions that tell an interpreter how to transform data in order to produce an output.
- When we want to use our model, we just load it into memory and execute it using the **TensorFlow interpreter**.
- TensorFlow's interpreter is designed to run models on powerful desktop computers and servers
- TensorFlow Lite Converter. The converter can also apply special optimizations aimed at reducing the size of the model and helping it run faster, often without sacrificing performance.

6. Run inference

```
private fun classifyDrawing() {  
    val bitmap = drawView?.getBitmap()  
  
    if ((bitmap != null) && (digitClassifier.isInitialized)) {  
        digitClassifier  
            .classifyAsync(bitmap)  
            .addOnSuccessListener { resultText -> predictedTextView?.text = resultText }  
            .addOnFailureListener { e ->  
                predictedTextView?.text = "classification error"  
                Log.e(TAG, "Error classifying drawing.", e)  
            }  
    }  
}
```

6. Run inference

STEPS TO USE THE TFLITE MODEL IN OUR APP

1. Load the TFLITE model from the asset folder – `loadModelFile()`
1. Initialize the interpreter - `initializeInterpreter()`
2. Prepare the data input in a byte buffer array
3. Run inference: `interpreter.run(input, output)`

See provided app for function implementation details

TensorFlow Lite

EE P 523, Lecture 5

TensorFlow *Lite*

TensorFlow: end-to-end open-source platform for machine learning.

TensorFlow Lite: lightweight solution for mobile and embedded devices.

- It enables on-device machine learning inference with low latency and a small binary size.
- TensorFlow Lite also supports hardware acceleration with the Android Neural Networks API.

TensorFlow *Lite*

Why do we need a new mobile-specific library?

The next wave of machine learning applications will have significant processing on mobile and embedded devices.

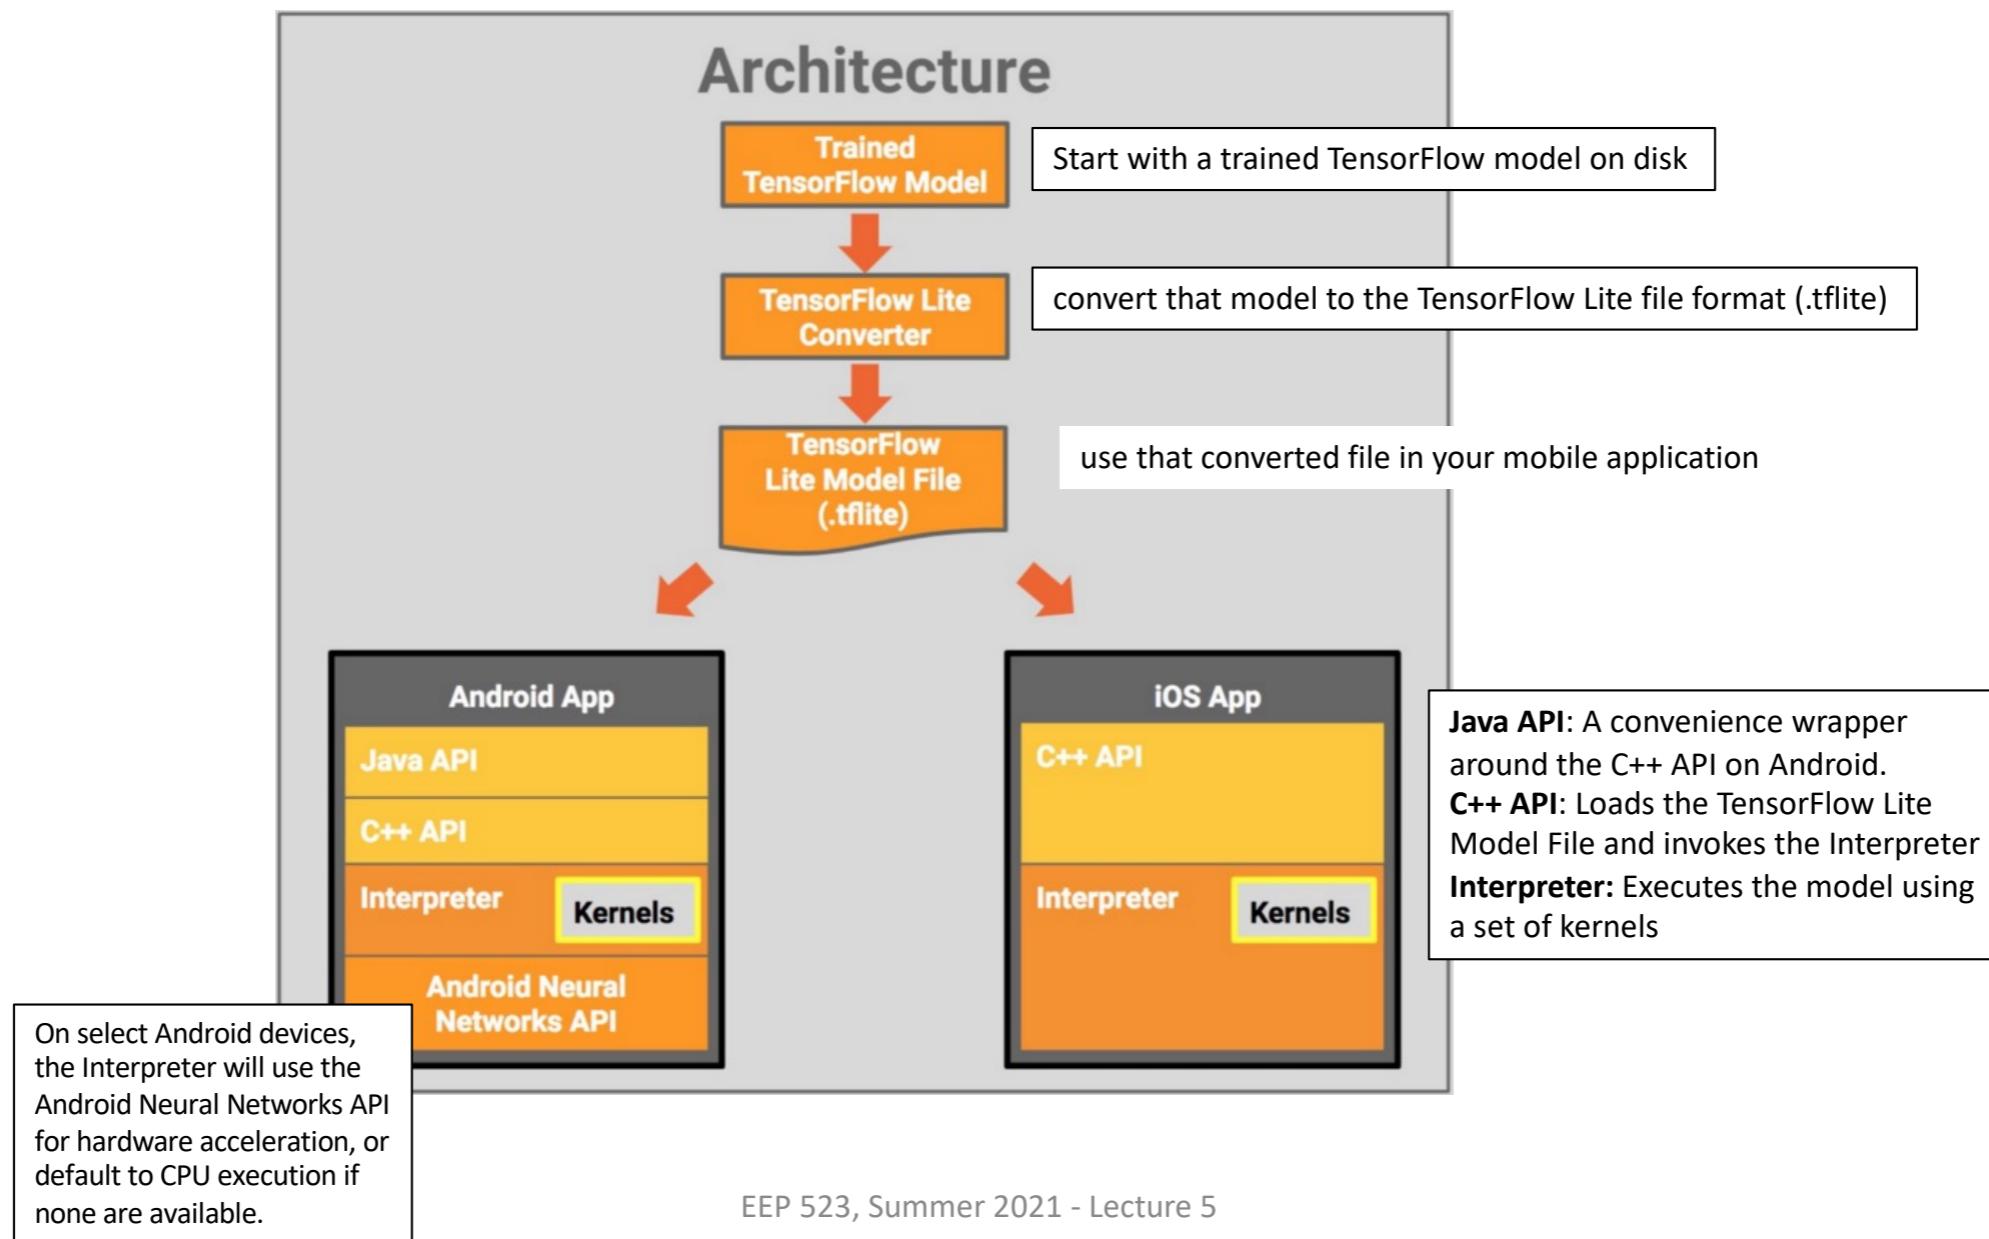
- Machine Learning is changing the computing paradigm.
- Consumer expectations are also trending toward natural, human-like interactions with their devices, driven by the camera and voice interaction models.

TensorFlow Lite

Why do we need a new mobile-specific library?

- Innovation at the *silicon layer* is enabling new possibilities for hardware acceleration, and frameworks such as the Android Neural Networks API make it easy to leverage these.
- Recent advances in real-time computer-vision and spoken language understanding have led to mobile-optimized benchmark models being open sourced (e.g., MobileNets, SqueezeNet).
- Widely-available smart appliances create new possibilities for on-device intelligence.
- Interest in stronger user data privacy paradigms where user data does not need to leave the mobile device.
- Ability to serve 'offline' use cases, where the device does not need to be connected to a network.

TensorFlow Lite



Why is it Called TensorFlow?

- **TensorFlow** is called 'TensorFlow' because it handles the flow (node/mathematical operation) of tensors, which are data structures that you can think of as multi-dimensional arrays.
- Tensors are represented as n-dimensional arrays of base datatypes, such as a string or integer -- they provide a way to generalize vectors and matrices to higher dimensions.
- The ``shape`` of a Tensor defines its number of dimensions, and the size of each dimension.
- The ``rank`` of a Tensor provides the number of dimensions (n-dimensions) -- you can also think of this as the Tensor's order or degree. Let's first look at 0-d Tensors, of which a scalar is an example:

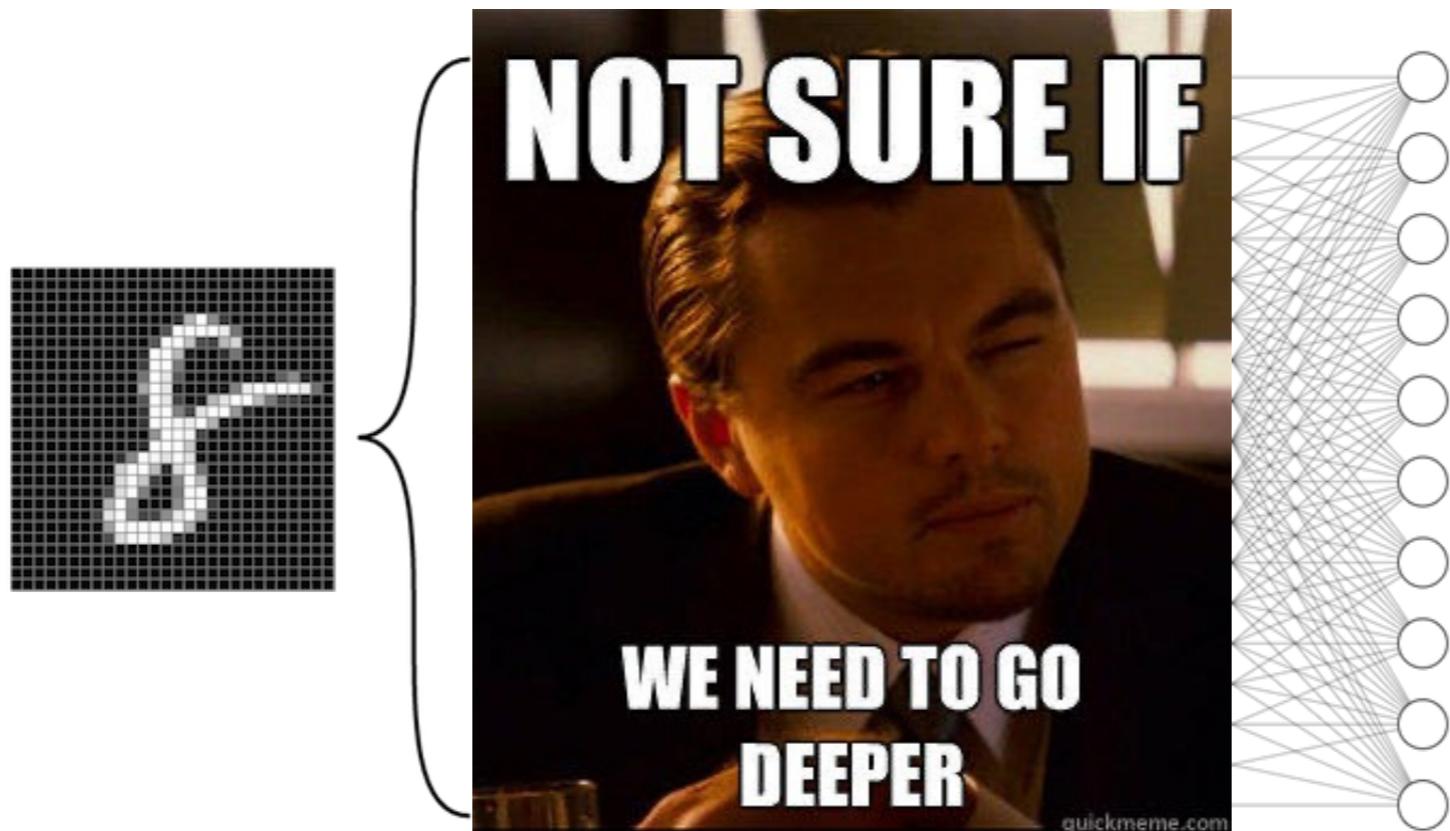
6. Run inference

- Since this is the part where our model meets our application code
- we need to write some code **that takes raw input data from our sensors and transforms it into the same form that our model was trained on.**
- We then pass this transformed data into our model and run inference.

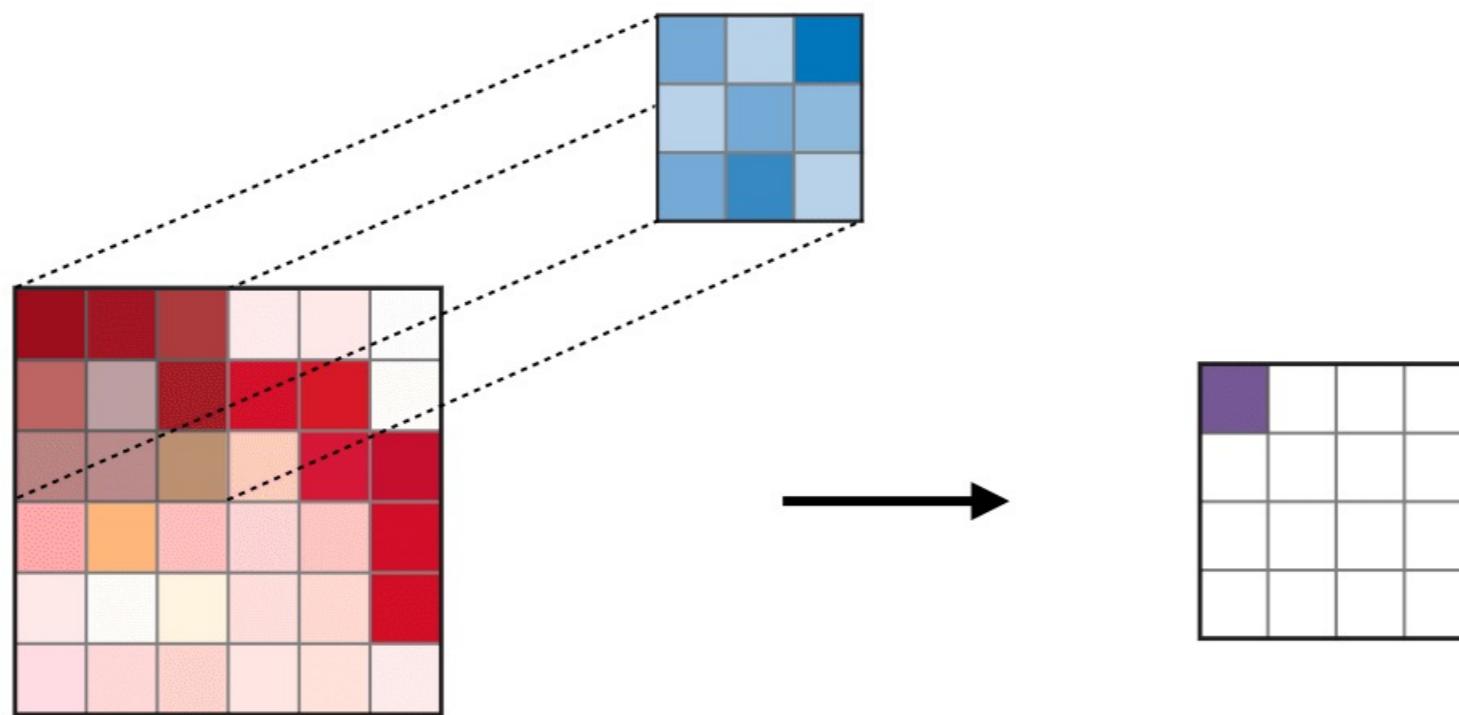
Improving Our Model

EE P 523, Lecture 5

How Do We improve the model?



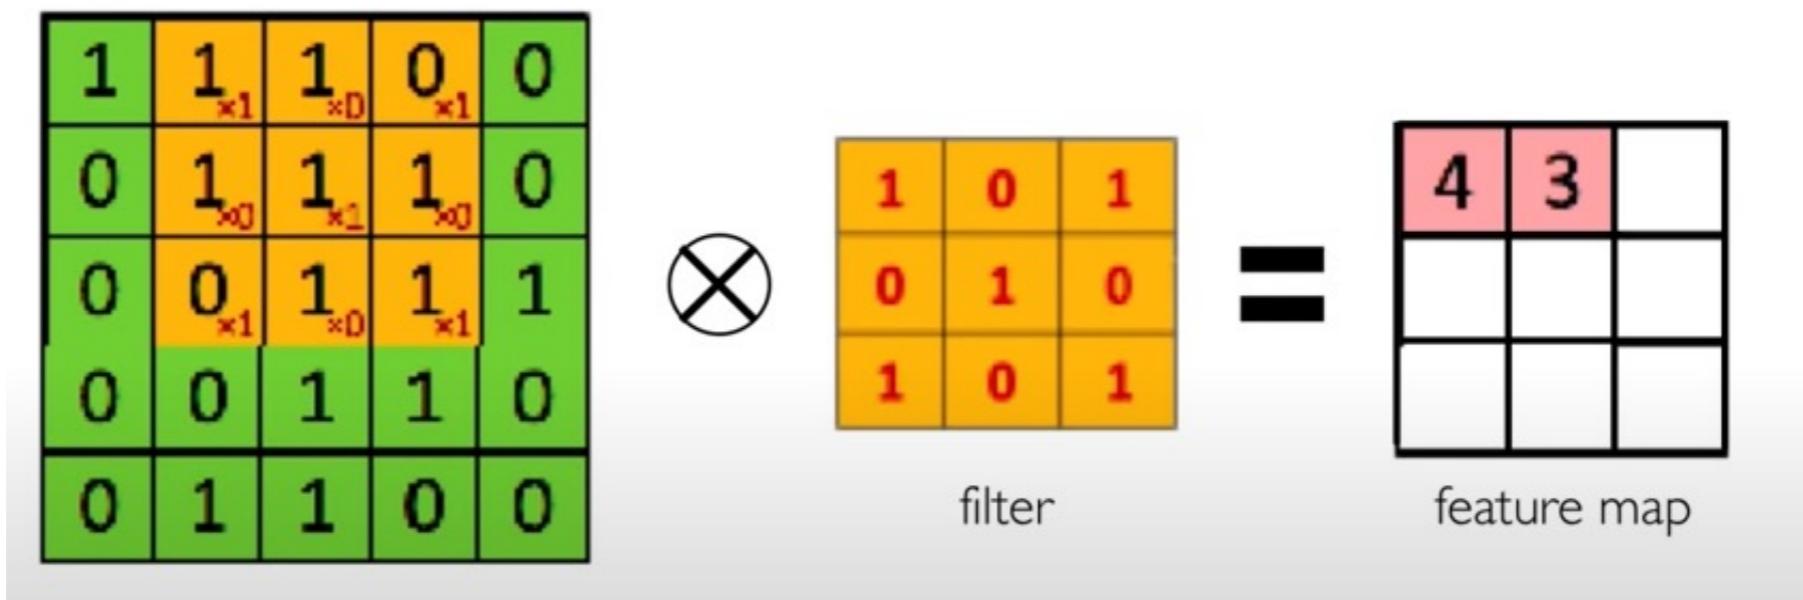
Convolution Operation



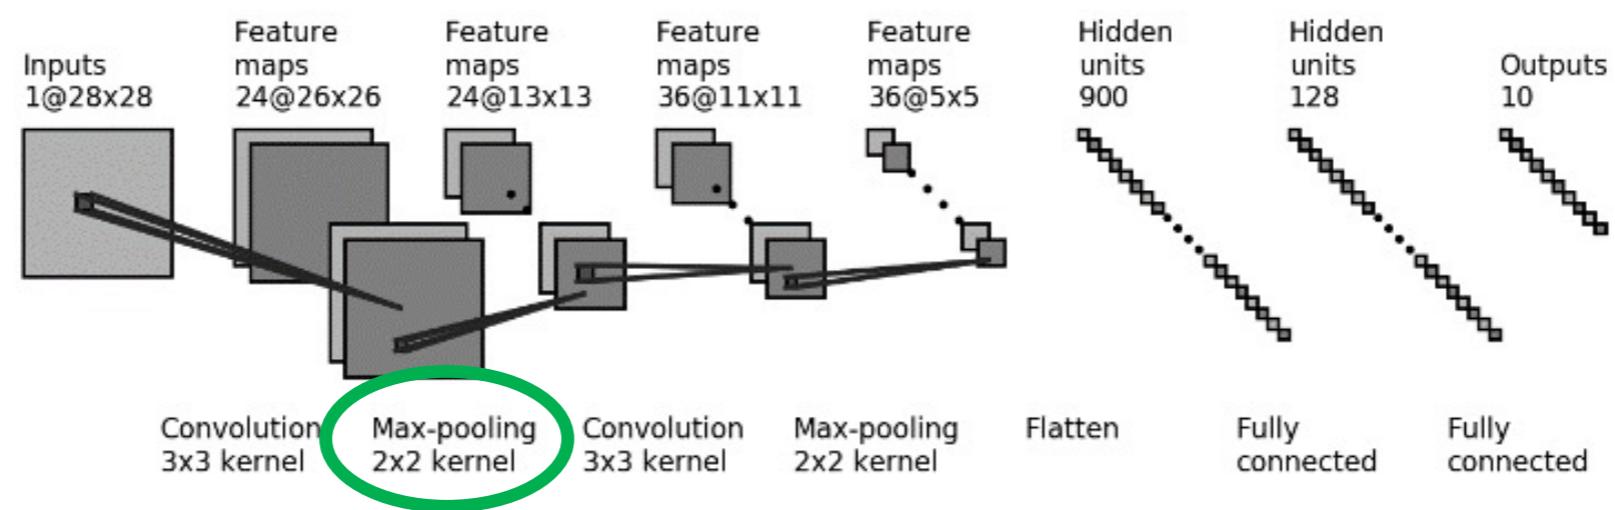
- Convolutional layers are used for **spotting 2D patterns in input images**.
- Each filter is a rectangular array of values that is moved as a sliding window across the input, and the output image is a representation of how closely

Convolution Operation

- 1)Slide a 3x3 filter
- 2)Element-wise multiply
- 3)Add the numbers

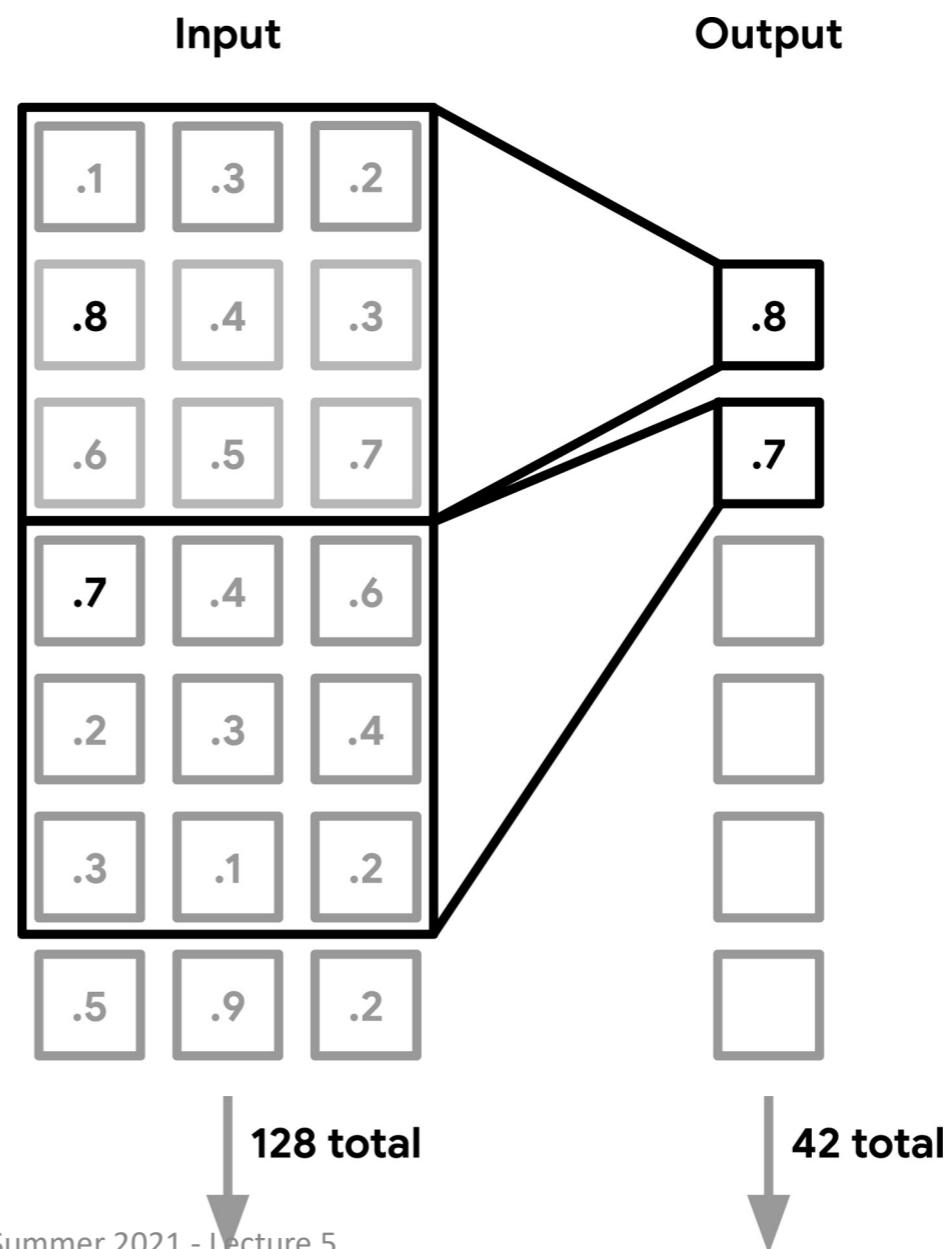


Convolutional Neural Network for digits recognition

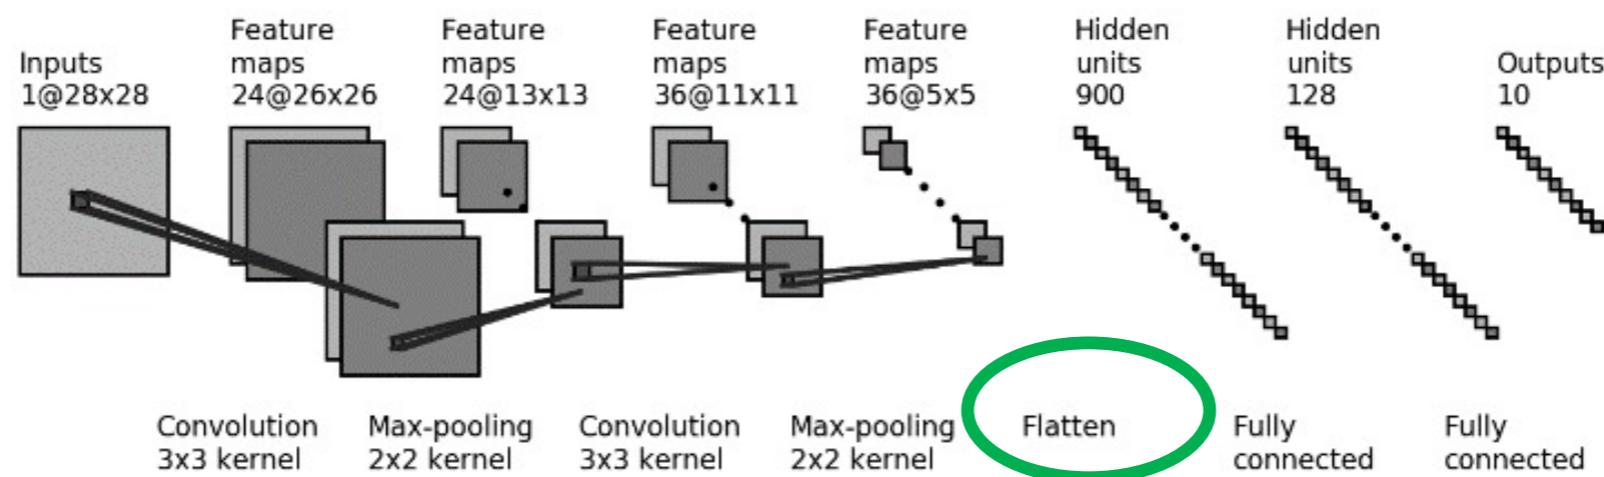


Max-pooling

- The pooling layer (POOL) is a **downsampling operation**, typically applied after a convolution layer, which does some spatial invariance.
- In particular, max and average pooling are special kinds of pooling where the maximum and average value is taken, respectively



Convolutional Neural Network for digits recognition



- The **Flatten layer** is used to **transform a multidimensional tensor into one with a single dimension**.
- In this case, our (36, 5, 5) tensor is squished down into a single dimension with shape (900).

Include the model into our Android App

The screenshot shows the Android Studio interface. On the left is the Project Navigational Drawer with sections for Project, Resource Manager, and Gradle Scripts. The Project section shows a tree view of the app module: app (selected), manifests, java, java (generated), assets (circled in red), res, and Gradle Scripts. In the main editor area, the build.gradle (App) file is open. The code is as follows:

```
include ':app'

dependencies {
    implementation 'org.tensorflow:tensorflow-lite:0.0.0-nightly'
}
```

A green callout box points to the assets folder in the project tree with the text "Include in the build.gradle(App)".

```
android{
    aaptOptions {
        noCompress "tflite"
    }
}

dependencies {
    implementation 'com.google.android.gms:play-services-tasks:17.0.0'
    implementation 'org.tensorflow:tensorflow-lite:0.0.0-nightly'
}
```

HOMEWORK HINTS



Assignment: Face Detection

In your main activity:

1. Create the Face Detector

```
val detector = FaceDetector.Builder(applicationContext)
    .setTrackingEnabled(false)
    .setLandmarkType(FaceDetector.ALL_LANDMARKS)
    .setClassificationType(FaceDetector.ALL_CLASSIFICATIONS)
    .build()
```

2. Create a frame from the bitmap and run the face detection on the frame

```
val frame = Frame.Builder().setBitmap(bitmap).build()
val faces = detector.detect(frame)
```

Assignment: Face Detection

In you View Class

- Receive the faces and the bitmap from the main activity
- Process the bitmap based on the landmarks
- Paint the result on a Canvas

```
class myViewClass(context: Context, attrs: AttributeSet) : View(context, attrs)
{
    private var mBitmap: Bitmap? = null
    private var mFaces: SparseArray<Face>? = null
    . . . . .
```

Assignment: Face Detection

In you View Class

- Scale the Bitmap to the size of the Canvas

```
val viewWidth = canvas.width.toDouble()
val viewHeight = canvas.height.toDouble()
val imageWidth = mBitmap!!.width.toDouble()
val imageHeight = mBitmap!!.height.toDouble()
val scale = Math.min(viewWidth / imageWidth, viewHeight / imageHeight)
```

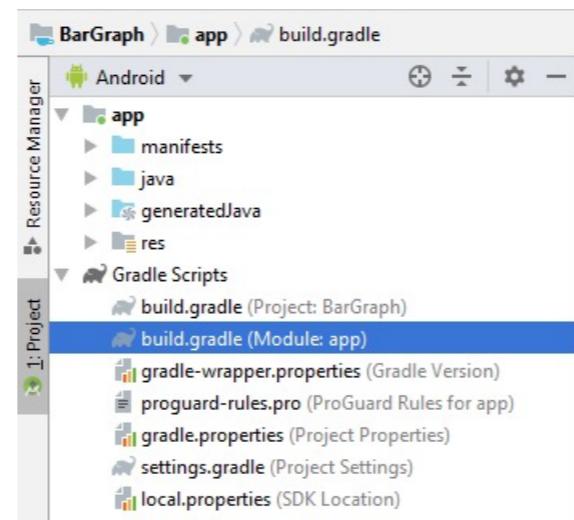
Assignment: Face Detection

In you View Class

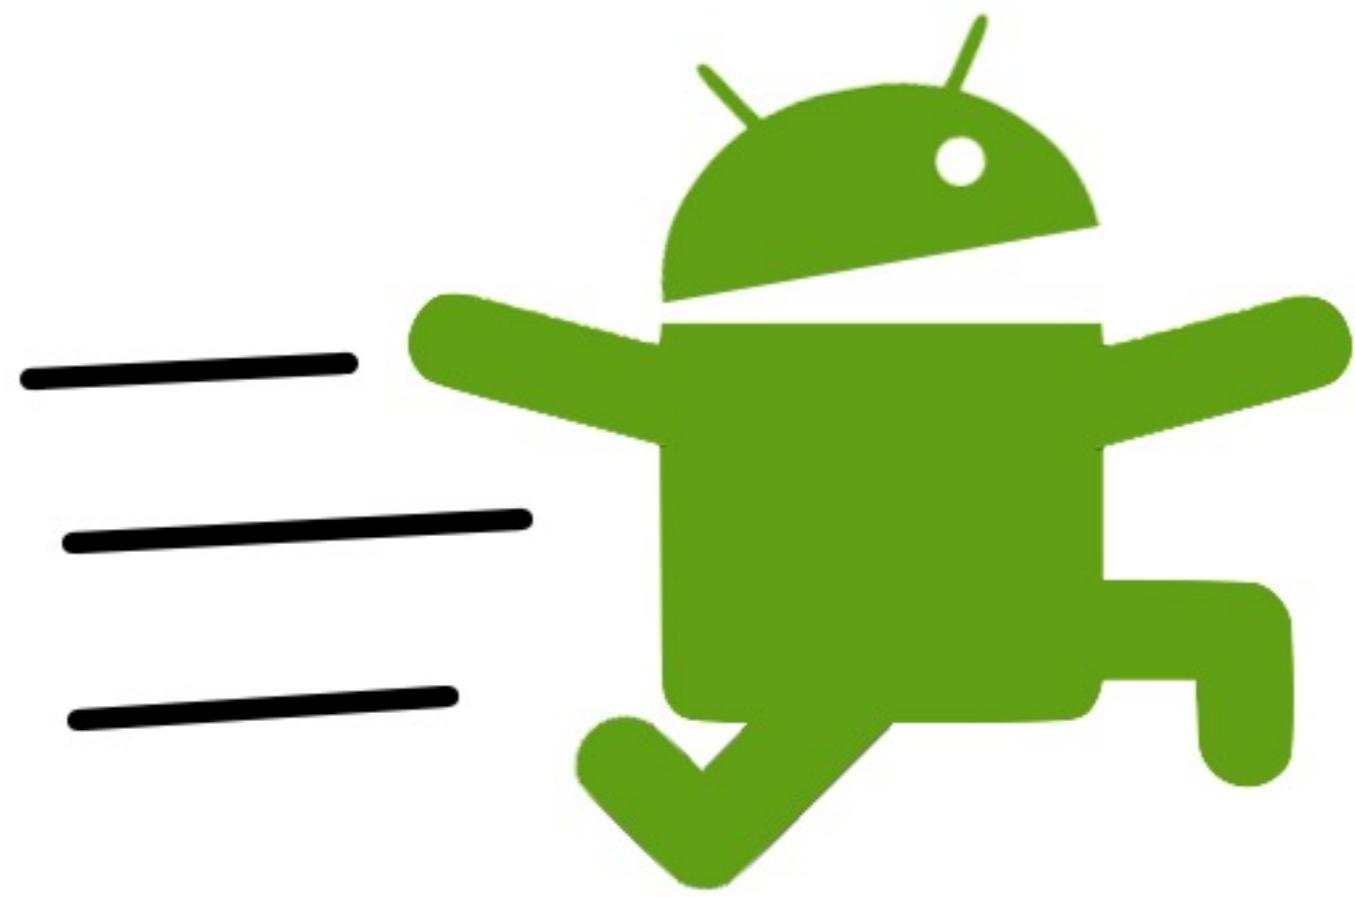
- Iterate over the faces, and recover the landmarks

```
for (i in 0 until mFaces!!.size()) {  
    val face = mFaces!!.valueAt(i)  
    var smileProb = face.isSmilingProbability  
    for (landmark in face.landmarks) {  
        val cx = (landmark.position.x * scale).toInt()  
        val cy = (landmark.position.y * scale).toInt()  
        //Example of plotting a circle in the landmark positions  
        canvas.drawCircle(cx.toFloat(), cy.toFloat(), 10f, paint)  
    }  
}
```

Assignment: Face Detection



```
dependencies {
    ...
    implementation 'com.google.android.gms:play-services-vision:9.4.0'
    ...
}
```



Your Questions

